

Rank Before You Merge: Native-Unit Selection for Visual Token Compression

Zhengxing Yan¹, Hao Sun², Chen Guo³, and Jianpeng Hu^{1,*}

¹Shanghai University of Engineering Science

²Nanjing University of Science and Technology

³East China Normal University

super_fw@163.com; sh1076849939@163.com

guochenjkhjk@qq.com; mr@sues.edu.cn

*Corresponding author

Abstract

Visual token compression depends on both the retained budget and the representation used to select it. We study this dependence in vision-language models with native spatial mergers. Rank-Before-Merge (RBM) scores complete native units using pre-merger feature norms and keeps their identities aligned with the model’s visual streams, without training or changing merger weights. At 25% retention, RBM exceeds post-merger norm selection on all nine text-dense model–benchmark pairs across Qwen3-VL, Qwen2.5-VL, and InternVL3. These operational comparisons include model-specific feature taps and allocation rules. To reduce the feature-depth confound, a separate full-split Qwen3-VL final-input control scores the main merger’s own input: it improves TextVQA by 27.68 percentage points, DocVQA by 4.59 points, and OCRBench by 235/1000, but reduces GQA by 5.64 points. The post score still aggregates main and auxiliary outputs, so this control is not a pure intervention on merger order. Rank disagreement and edge-energy associations characterize the selected evidence; a ranking-swap check verifies path equivalence for a fixed kept set. Comparisons with FastV show complementary task preferences rather than uniform superiority. Code: github.com/Code-Yan-ZX/qrbm-vlm.

Introduction

High-resolution images produce long visual sequences in modern vision-language models (VLMs). A common architectural response is a native spatial merger: neighboring patch features are converted into fewer language-model inputs. Additional compression must then decide which of these spatial units deserve a smaller token budget. This decision depends on the scoring representation. A large norm before a learned projection need not remain large afterward, and a downstream selector cannot recover evidence from a unit it discards.

We investigate this choice using a deliberately simple, query-blind norm score. Rank-Before-Merge (RBM) selects complete native units from pre-merger features while retaining native merger weights and downstream interface dimensions. It requires neither training nor a second learned compression module. Its purpose is to expose a useful selection signal while respecting the model’s existing spatial topology.

Two questions must be separated. First, does an operational RBM implementation preserve accuracy at the same final token budget as post-merger norm selection? Second, does a later pre-merger tap retain the same direction of effect after removing




Figure 1: Native-unit selection schematic. Selected identities stay aligned with the surviving visual streams and interface. The illustrated pre-merger location is conceptual: operational Qwen3-VL RBM scores the layer-8 deepstack input, whereas the final-input control scores the main merger’s final-block input. The diagram’s ρ denotes retention, written κ in the text.

the early-feature advantage? Our three-model comparison addresses the first question. A separate Qwen3-VL experiment scores the main merger’s final-block input while holding sample identities and token counts fixed. It reduces, but does not eliminate, implementation differences because the post score concatenates main and auxiliary-stream outputs. This distinction matters because operational Qwen3-VL RBM scores an earlier tap and prunes multiple visual streams.

At 25% retention, operational RBM improves all nine text-dense comparisons against Post-L2 across three model families. The final-input control also favors its pre-merger score on TextVQA, DocVQA, and OCRBench, but the DocVQA gain is substantially smaller. GQA reverses the direction. Thus neither an architecture-independent preference for earlier features nor a universal advantage over stronger selectors follows from the results.

Our contributions are a native-unit compression rule compatible with deployed merger interfaces; a full-split evaluation separating operational performance from a stricter final-input control; and diagnostics and attention-based comparisons that identify the scope of the observed advantage. We interpret visual-token retention as a task compression budget, without equating it to a bit rate or to total inference cost.

Related Work

Visual-token reduction can operate inside the language model through attention or progressive pruning [1, 2, 3, 4], or at the encoder output through [CLS] attention, graph propagation, fitted predictors, and multiple cues [5, 6, 7, 8, 9]. Unified systems combine saliency with diversity or text guidance [10, 11, 12]. Their signals and pruning stages differ from RBM’s image-only norm score, and equal final token count need not imply equal computation.

Merging changes the representation rather than only deleting tokens. ToMe, VisionZip, and AdaptMerge use matching, dominant-token preservation, and language guidance [13, 14, 15]; TokenPacker learns a compact projector, while GlimpsePrune, VScan, and AnchorPrune adapt reduction to inputs or model families [16, 17, 18, 19]. These methods motivate treating the selector’s input representation as part of the compression rule.

Recent selectors move reduction earlier or vary it across tokens and layers [20, 21, 22, 23, 24]; optimize coverage, reconstruction, or sensitivity [25, 26, 27, 28, 29]; and model semantics, information flow, or head importance [30, 31, 32, 33]. Layerwise, object-centric, local/global, and hybrid designs preserve evidence in other ways [34, 35, 36, 37, 38, 39]. Document pruning and provenance audits expose OCR-specific failures [40, 41], while LMMS-Eval and EffiVLM-Bench support broader evaluation [42, 43].

We focus on architectures with native spatial compression: Qwen2.5-VL, Qwen3-VL, and InternVL3 [44, 45, 46]. RBM retains their learned native operators and selects complete input units. Post-L2 is a diagnostic baseline from the same feature-magnitude family; FastV supplies a stronger attention-based comparison. The resulting evidence concerns these implementations and workloads, rather than a claim that downstream selectors generally assume rank preservation or that RBM ranks the full design space above.

Native-Unit Compression

Selection representation and token budget

Let an image contain N native spatial units, each grouping four patch vectors. At retention κ , retain

$$k = \max\{1, \text{round}(\kappa N)\}. \quad (1)$$

For features $h_{u,j}^{(t)}$ at tap t , RBM uses

$$a_u^{(t)} = \frac{1}{4} \sum_{j=1}^4 \|h_{u,j}^{(t)}\|_2, \quad K_{\text{RBM}} = \text{TopK}_k\{a_u^{(t)}\}. \quad (2)$$

Post-L2 instead scores a returned visual vector z_u :

$$b_u = \|z_u\|_2, \quad K_{\text{post}} = \text{TopK}_k\{b_u\}. \quad (3)$$

Both rules use magnitude, but they evaluate different representations; they are not identical scoring functions applied to identical vectors. In Qwen3-VL, z_u concatenates the main and three deepstack outputs, whereas the other tested paths use their model-specific post-projector vectors.

In the idealized single-stream case with $h^{(t)}$ at the merger input, $z_u = M(h_{u,1:4}^{(t)})$. For a per-unit merger and a fixed kept set, selection can be performed before or after M without changing the retained outputs. In contrast, the ranking decisions coincide only when their top- k sets agree. A strictly increasing map $b_u = g(a_u)$ is sufficient for agreement at all budgets, but is not a property we assume of learned mergers. We

quantify set agreement with Jaccard overlap, $|K_{\text{pre}} \cap K_{\text{post}}| / |K_{\text{pre}} \cup K_{\text{post}}|$. Low overlap alone predicts neither which set is more useful nor the size of an accuracy difference. Qwen3-VL’s concatenated post vector departs from this single-stream idealization, so its final-input experiment is a stricter control rather than exact causal isolation.

Model-specific implementation

RBM retains complete units and keeps selected features aligned with the positions and auxiliary streams expected by each adapter. Model weights, merger weights, interface dimensions, and decoding are unchanged. The implementations follow their backend’s compressed-position convention; this does not imply that every backend gathers the selected units’ original spatial coordinates. L2 scoring costs $O(Nd)$ for feature dimension d ; top- k can be implemented in $O(N \log k)$ with a heap. The method adds no learned parameters or language-model attention pass.

For Qwen3-VL, operational RBM scores the first deepstack input at ViT layer 8 and applies one kept set to the main and three deepstack mergers. Post-L2 scores the concatenated returned vector. Consequently, their operational comparison changes feature depth, score aggregation, and where auxiliary streams are pruned. Qwen2.5-VL scores the main merger input, but its available pre path keeps units in window order while the post path follows the native output restoration. InternVL3 scores the concatenated four-patch vector after parameter-free pixel shuffle and before its learned projector; RBM allocates a quota per tile, whereas Post-L2 pools each image’s tiles before top- k . These results are therefore operational accuracy comparisons, not isolated cross-model merger effects.

The separate Qwen3-VL *Pre-final* control scores final ViT-block features at the main merger input. Pre-final and Post-L2 both execute the deepstack mergers in full and select corresponding rows afterward. Their configuration, sample identities, and per-image token counts match. This removes the early-tap difference, but the pre score still uses four input patch norms while Post-L2 uses a concatenated main-plus-deepstack vector. We therefore call it a final-input control, not a pure stage intervention. It is also not an efficiency measurement of physically skipping merger work, and RBM–Pre-final differences are not pure feature-depth effects.

Evaluation Protocol

We evaluate Qwen3-VL-8B, Qwen2.5-VL-7B, and InternVL3-8B with bf16 weights and greedy decoding. The main matrix and final-input control use vLLM 0.19 V1 [47]. FastV and the hybrid probe use HuggingFace eager attention. We compare methods within each backend; small checks matching native outputs on 8/8 examples and HF/vLLM outputs on 16/16 examples do not establish full-split backend equivalence.

The full evaluation sets contain 5,000 TextVQA examples, 5,349 DocVQA examples, 1,000 OCRBench examples, and 12,578 GQA examples. Metrics are VQA accuracy, ANLS, OCRBench score out of 1,000, and normalized exact match, respectively [48, 49, 50, 51]. Scoring follows the corresponding lmms-eval rules [42]. The three former benchmarks are designated text-dense; GQA provides an object-centered question-answering contrast.

Table 1: Full-split operational results at 25% retention. None is uncompressed. Qwen3-VL RBM uses the early tap; this table does not isolate scoring stage. Deltas use unrounded scores and are percentage points except for OCRBench (/1000).

Model	Benchmark	None	RBM	Post-L2	Δ
Qwen3-VL	TextVQA	.844	.605	.222	+38.4
	DocVQA	.956	.481	.238	+24.3
	OCRBench	760	547	184	+363
	GQA	.616	.449	.477	-2.8
Qwen2.5-VL	TextVQA	.862	.702	.442	+26.1
	DocVQA	.949	.636	.526	+11.0
	OCRBench	817	480	182	+298
	GQA	.604	.559	.585	-2.6
InternVL3	TextVQA	.8338	.7890	.4148	+37.4
	DocVQA	.9221	.7284	.3820	+34.6
	OCRBench	852	753	321	+432
	GQA	.6293	.5993	.6031	-0.4

All reported compression comparisons retain 25% of native units and match per-image final visual-token counts within each experiment. Qwen3-VL final-input runs use native image resolution, at most 32 generated tokens, and eager vLLM execution. TextVQA, OCRBench, and GQA use a maximum sequence length of 8,192 and up to eight concurrent sequences; DocVQA uses 32,768 and up to four. Both control arms complete every evaluation sample. The reused uncompressed Qwen3-VL anchors score two TextVQA and 18 OCRBench context overruns as zero. Qwen2.5-VL’s operational OCRBench pair uses a shared four-million-pixel cap. InternVL3 DocVQA also uses a four-million-pixel cap. Absolute scores across model families are therefore not controlled model rankings.

For the final-input control, 95% confidence intervals use 20,000 paired-bootstrap resamples of per-example official scores. We report these intervals separately from the exploratory hybrid probe, for which we make no significance claim. Full-split scores, fixed-size probes, and skipped-example handling are distinguished below.

Results

Operational accuracy across model families

Table 1 compares the deployed RBM design with Post-L2. RBM improves all nine text-dense model–benchmark pairs: 11.0–38.4 percentage points on TextVQA and DocVQA, and 298–432 points on OCRBench. It nevertheless remains lossy relative to the uncompressed models. For example, Qwen3-VL DocVQA falls from .956 to .481, despite its advantage over Post-L2. A favorable relative comparison is not equivalent to preserving the original model’s accuracy.

GQA favors Post-L2 in all three point estimates. The small InternVL3 difference has an interval crossing zero in the paired analysis. The Qwen results give clearer evidence of a task boundary: pre-merger magnitude is useful for the tested text-dense

Table 2: Qwen3-VL final-input control, complete full splits. Pre-final scores four input patches; Post-L2 scores the concatenated main-plus-deepstack output. Deltas and 95% paired-bootstrap intervals are percentage points, except OCRBench points.

Benchmark	Pre-final	Post-L2	Δ [95% CI]
TextVQA	.4985	.2217	+27.68 [26.21, 29.19]
DocVQA	.2836	.2377	+4.59 [3.43, 5.78]
OCRBench	419	184	+235 [199, 271]
GQA	.4207	.4771	−5.64 [−6.42, −4.86]




Figure 2: Full-split Qwen3-VL accuracy at 25% retention. Early-tap RBM is the operational design; Pre-final and Post-L2 form the stricter final-input comparison. OCRBench is normalized by 1,000. The RBM–Pre-final difference also includes auxiliary-stream handling and is not a feature-depth-only intervention.

workloads, but is not a general-purpose saliency criterion.

Final-input control

Table 2 and Fig. 2 separate the Qwen3-VL final-input contrast from the operational early-tap design. Pre-final improves TextVQA by 27.68 points and OCRBench by 235/1000. DocVQA improves by 4.59 points, compared with the 24.3-point operational contrast. These are distinct effects; attributing the latter entirely to the main merger boundary would overstate the evidence. Conversely, GQA decreases by 5.64 points in the final-input comparison. The full-split control therefore shows that a later pre-merger representation retains a workload-dependent effect, while its residual score-aggregation difference limits a merger-only attribution.

Comparison with attention-based selection

A same-harness full-OCRBench comparison evaluates RBM and FastV-k3 at a shared four-million-pixel cap and a 25% final budget (Table 3). Both methods use native positions, greedy decoding, and at most 32 generated tokens. The Qwen3-VL and Qwen2.5-VL runs encounter 79 and 88 common context or memory failures, respectively. These exact shared IDs receive zero in the scores out of 1,000; the completed paired sets contain 921 and 912 examples. RBM leads by 146 and 87 points. These

Table 3: Same-harness HF eager OCRBench comparison, 1,000 attempted examples per model. Common failures are scored zero; both methods retain the same final token count.

Model	RBM	FastV-k3	Δ (/1000)	Completed
Qwen3-VL	559	413	+146	921
Qwen2.5-VL	382	295	+87	912

HF scores are a separate experiment and must not be substituted into the vLLM matrix.

Diagnostics and Interpretation

The two operational Qwen3-VL rankings differ substantially. On diagnostic DocVQA, TextVQA, and GQA samples, Spearman correlations are .137, .332, and .360, and top-25% Jaccard overlaps are .180, .243, and .278. Representation choice therefore changes the selected set, although the different feature depths prevent merger-only attribution.

Rank shifts also associate with local edge energy: post-minus-pre rank correlations with Sobel energy are +.439, +.155, and +.036 in the same task order. On DocVQA, RBM-only versus Post-L2-only units have mean energy .641 versus .124, and 92% versus 35% exceed their image’s median. Because Sobel energy also responds to borders, texture, and contours, this establishes an edge association rather than text-specific or causal demotion.

A ranking-swap check fixes the post-merger computation but substitutes RBM’s kept identities. On the verified Qwen3-VL path, kept-set overlap is 1.000, retained representations are byte-identical, and predictions agree. This verifies path equivalence for a fixed set, not that norm or edge energy identifies causally relevant text. Qwen2.5-VL retains an ordering confound, while InternVL3 provides accuracy-level replication without byte-exact attribution.

Cost and Limitations

The retained-unit fraction k/N measures the final visual sequence budget, not an encoded bit rate or total inference cost. FastV processes all candidates through early language-model layers, while RBM does not eliminate vision-transformer work before its scoring tap; Pre-final deliberately executes all merger streams. Resolution, output length, batching, and backend also affect wall-clock performance.

On one A40, a Qwen3-VL TextVQA batch of 200 examples reaches 6.91 requests/s with RBM at 25% retention, versus 4.11 uncompressed and 6.89 with Post-L2. The vLLM run uses native resolution, at most eight concurrent sequences, greedy decoding capped at 32 tokens, and one untimed warmup; timing excludes engine construction and initial image loading. A separate five-run, one-example proxy averages .23 seconds for RBM and .36 seconds uncompressed. This supports an end-to-end gain for that workload, while the equal RBM/Post-L2 throughput shows no selector-specific speed advantage; the proxy is not a serving-latency distribution.

Our strongest controlled conclusion is the full-split Qwen3-VL final-input comparison, subject to its score-aggregation limitation. The three-model matrix extends the pattern across native merger designs without establishing a universal mechanism. Post-L2 is diagnostic, and FastV shows that RBM is not uniformly strongest. Shared loading failures also limit the HF OCRBench comparison. Broader selectors and direct text-region interventions remain necessary for general competitiveness or semantic claims.

Conclusion

Selecting native visual units from pre-merger features is a useful training-free compression design for the tested text-dense workloads. Separating the deployed RBM implementation from a final-input control shows that the useful signal persists at a later pre-merger tap, while score aggregation, feature depth, and stream handling also matter to operational performance. The GQA reversal and FastV comparisons define a clear boundary: selection representation should be evaluated against the intended workload, alongside the token budget and the model’s native interface.

References

- [1] L. Chen *et al.*, “An image is worth 1/2 tokens after layer 2: Plug-and-play inference acceleration for large vision-language models,” in *Eur. Conf. Comput. Vis. (ECCV)*, 2024.
- [2] L. Xing *et al.*, “Pyramiddrop: Accelerating your large vision-language models via pyramid visual redundancy reduction,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2025.
- [3] Q. Zhang *et al.*, “Beyond text-visual attention: Exploiting visual cues for effective token pruning in vlms,” 2024.
- [4] Y. Zhang *et al.*, “Sparsevlm: Visual token sparsification for efficient vision-language model inference,” in *Proc. Int. Conf. Mach. Learn. (ICML)*, ser. Proceedings of Machine Learning Research, vol. 267. PMLR, 2025.
- [5] A. Wang *et al.*, “[CLS] token tells everything needed for training-free efficient MLLMs,” 2024.
- [6] Y. Shang *et al.*, “Llava-prumerge: Adaptive token reduction for efficient large multi-modal models,” in *Proc. IEEE/CVF Int. Conf. Comput. Vis. (ICCV)*, 2025.
- [7] W. Ye *et al.*, “Fit and prune: Fast and training-free visual token pruning for multi-modal large language models,” in *Proc. AAAI Conf. Artif. Intell. (AAAI)*, 2025.
- [8] Y. Jiang *et al.*, “What kind of visual tokens do we need? training-free visual token pruning for multi-modal large language models from the perspective of graph,” 2025.
- [9] B. Luan *et al.*, “Multi-cue adaptive visual token pruning for large vision-language models,” 2025.
- [10] H. Yu *et al.*, “VisionTrim: Unified vision token compression for training-free MLLM acceleration,” in *Int. Conf. Learn. Represent. (ICLR)*, 2026.
- [11] C. Baek *et al.*, “AgilePruner: An empirical study of attention and diversity for adaptive visual token pruning in large vision-language models,” in *Int. Conf. Learn. Represent. (ICLR)*, 2026.
- [12] Z. Fang *et al.*, “Prune redundancy, preserve essence: Vision token compression in VLMs via synergistic importance-diversity,” in *Int. Conf. Learn. Represent. (ICLR)*, 2026.

- [13] D. Bolya *et al.*, “Token merging: Your ViT but faster,” in *Int. Conf. Learn. Represent. (ICLR)*, 2023.
- [14] S. Yang *et al.*, “Visionzip: Longer is better but not necessary in vision language models,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2025.
- [15] Z. Islam and M. Rochan, “AdaptMerge: Inference time adaptive visual and language-guided token merging for efficient large multimodal models,” in *Findings Assoc. Comput. Linguistics: EMNLP*, 2025.
- [16] W. Li *et al.*, “TokenPacker: Efficient visual projector for multimodal LLM,” 2024.
- [17] Q.-S. Zeng *et al.*, “A glimpse to compress: Dynamic visual token pruning for large vision-language models,” *IEEE Trans. Circuits Syst. Video Technol.*, pp. 1–1, 2026.
- [18] C. Zhang *et al.*, “Vscan: Rethinking visual token reduction for efficient large vision-language models,” 2025.
- [19] K. Oh and B. Kim, “AnchorPrune: Relevance-anchored contextual expansion for visual token pruning,” in *Eur. Conf. Comput. Vis. (ECCV)*, 2026.
- [20] T. Gao *et al.*, “QuietPrune: Query-guided early token pruning for vision-language models,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2026, pp. 3553–3562.
- [21] N. Wang *et al.*, “PixelPrune: Pixel-level adaptive visual token reduction via predictive coding,” 2026.
- [22] A. Li *et al.*, “TransPrune: Token transition pruning for efficient large vision-language model,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2026, pp. 39 529–39 538.
- [23] J. Chen *et al.*, “Variation-aware vision token dropping for faster large vision-language models,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2026, pp. 3489–3499.
- [24] Z. Sun *et al.*, “Hi-Lo Prune: Look at what you’ll lose before pruning with hierarchical token selection,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2026, pp. 31 941–31 951.
- [25] Q. Ma *et al.*, “ApET: Approximation-error guided token compression for efficient VLMs,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2026.
- [26] T. Xu, H. Shi, and X. Gao, “SCoRe: Saliency-coverage reduction for vision token pruning in vision-language models,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2026.
- [27] C. Du *et al.*, “CoIn: Coverage and informativeness-guided token reduction for efficient large multimodal models,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2026, pp. 10 492–10 501.
- [28] Y. Kim *et al.*, “ZOO-Prune: Training-free token pruning via zeroth-order gradient estimation in vision-language models,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2026, pp. 39 572–39 582.
- [29] J. Lee, S. Wen, and D.-W. Choi, “Moving beyond diversity: Visual token pruning as subspace reconstruction for efficient VLMs,” 2026.
- [30] Y. Wang *et al.*, “Rethinking token reduction for large vision-language models,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2026, pp. 24 727–24 737.
- [31] J. Song *et al.*, “EvoComp: Learning visual token compression for multimodal large language models via semantic-guided evolutionary labeling,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2026.
- [32] G. Sun *et al.*, “IF-Prune: Information-flow guided token pruning for efficient vision-language models,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2026, pp. 3522–3531.

- [33] Q. Zhu *et al.*, “HAWK: Head importance-aware visual token pruning in multimodal models,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2026.
- [34] Y. Wang *et al.*, “When token pruning is worse than random: Understanding visual token information in VLLMs,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2026, pp. 31 910–31 919.
- [35] Y. Chen *et al.*, “One layer’s trash is another layer’s treasure: Adaptive layer-wise visual token selection in LVLMs,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2026.
- [36] X. Li *et al.*, “CRISP: Pre-LLM yet text-driven visual token pruning for efficient LVLM inference,” 2026.
- [37] J. Lei, G. Wang, and D.-H. Lee, “CORE: Compact object-centric REpresentations as a new paradigm for token merging in LVLMs,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2026.
- [38] J. Li *et al.*, “Token reduction via local and global contexts optimization for efficient video large language models,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2026.
- [39] J. Zhang *et al.*, “Hybrid token compression for vision-language models,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2026, pp. 31 889–31 899.
- [40] J. Choi *et al.*, “DocPrune: Efficient document question answering via background, question, and comprehension-aware token pruning,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2026, pp. 3543–3552.
- [41] F. Liu *et al.*, “Beyond accuracy: Auditing spatial provenance in visual token pruning for OCR-critical MLLM inference,” 2026.
- [42] K. Zhang *et al.*, “LMMs-Eval: Reality check on the evaluation of large multimodal models,” in *Findings Assoc. Comput. Linguistics: NAACL*, 2025, pp. 881–916.
- [43] Z. Wang *et al.*, “EffiVLM-BENCH: A comprehensive benchmark for evaluating training-free acceleration in large vision-language models,” in *Annu. Meeting Assoc. Comput. Linguistics (ACL)*, 2025.
- [44] S. Bai *et al.*, “Qwen2.5-vl technical report,” 2025.
- [45] —, “Qwen3-vl technical report,” 2025.
- [46] J. Zhu *et al.*, “Internvl3: Exploring advanced training and test-time recipes for open-source multimodal models,” 2025.
- [47] W. Kwon *et al.*, “Efficient memory management for large language model serving with PagedAttention,” in *Proc. 29th Symp. Oper. Syst. Princ. (SOSP)*, 2023.
- [48] A. Singh *et al.*, “Towards VQA models that can read,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2019.
- [49] M. Mathew *et al.*, “Docvqa: A dataset for vqa on document images,” in *Proc. IEEE/CVF Winter Conf. Appl. Comput. Vis. (WACV)*, 2021, pp. 2200–2209.
- [50] Y. Liu *et al.*, “OCRBench: On the hidden mystery of OCR in large multimodal models,” *Sci. China Inf. Sci.*, vol. 67, no. 22, p. 220102, 2024.
- [51] D. A. Hudson and C. D. Manning, “GQA: A new dataset for real-world visual reasoning and compositional question answering,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2019.