



pre - post (pp)

Qwen3-VL-8B Qwen2.5-VL-7B

TextVQA



DocVQA



Visual-token retention (deeper compression →)

n = 200 per benchmark; shaded region = deepest tested compression