

Rank Before You Merge: A Stage Law for Training-Free Visual Token Compression in Merger-Equipped Vision-Language Models

Zhengxing Yan

Abstract

Native spatial mergers reduce visual-token count before the language model, but whether they should precede token selection is under-controlled. We compare two iso-budget arms that score the same 2×2 units with query-blind L2 on merger-input (pre) or merger-output (post) features, holding model and decoding fixed. Across Qwen3-VL-8B, Qwen2.5-VL-7B, and InternVL3-8B, pre leads the matched L2-family post arm on all nine official TextVQA, DocVQA, and OCRBench cells at 25% retention, by +11.0 to +38.4 percentage points and +293 to +432 OCRBench points. A matched-depth control at the main-merger boundary retains the text-dense effect but detects no GQA difference.

On Qwen3-VL, the merger weakly preserves unit ranks and preferentially demotes document text-stroke units; swapping the pre ranking into the post path recovers pre accuracy with identical selected-unit identities. We turn this finding into **Rank-Before-Merge (RBM)**, a minimal, training-free stage change with no learned scorer or query input. The gains are against the matched L2-family post arm, not a universal method ranking: query-conditioned FastV remains stronger on several reference-grounded and object-centric settings, while RBM is the more robust OCR-oriented default. Six prespecified extensions do not improve over their stronger constituents. At 25% retention, compression is stage-neutral in speed, increasing throughput by 68% on Qwen3-VL and 1.8–2.5 $\times$ on InternVL3.

Keywords

visual token compression, vision-language models, training-free inference, multimodal foundation models, inference efficiency

1 Introduction

Multimodal inference cost is a visual-token problem. Visual-language models (VLMs) process images as hundreds to tens of thousands of visual tokens, making quadratic attention and key-value memory major inference costs. Visual-token compression is therefore central to efficient document understanding, scene-text QA, and OCR. In merger-equipped architectures such as Qwen-VL and InternVL3, 2×2 patch groups are first merged, while prior work focuses mainly on which tokens to keep and how to score them. The selection stage, before or after the native merger, is rarely isolated as a matched variable. Prior work notes Qwen’s native token reduction [3] and OCR degradation from post-encoder merging [13, 41]; we directly test the boundary by scoring the same units with query-blind L2 on merger-input (pre) or merger-output (post) features at the same per-image budget.

Minimal operationalization: Rank-Before-Merge (RBM). Score each merger-input unit, retain the top- κ , then merge only the survivors. RBM uses no attention weights, training, or query input. Because merging is per-unit, a fixed kept set receives the same native per-unit merger operation; byte identity is verified on

Qwen3-VL. RBM changes which units reach the merger. We use this rule to expose the stage variable, not to claim scorer novelty (§3.3).

What falls out of it — a stage law across models. Under iso-model, iso-budget, and a shared query-blind L2-family rule, pre leads post by 11.0–38.4 pp on official text-dense splits across two Qwen generations and by 34.6–37.4 pp plus 432 OCRBench points on InternVL3, whose merger design differs. A matched-depth pure-stage control retains the TextVQA/DocVQA effect but detects no GQA difference; the result is workload-conditional and widens with compression depth (§4).

Why the stage matters. On Qwen3-VL, the learned merger weakly preserves unit ranks (Spearman $\rho = 0.14\text{--}0.36$), preferentially demotes high-edge text-stroke units on documents, and yields to a ranking-swap control: the post path recovers pre accuracy with identical selected-unit sets (Jaccard = 1.000 in four model $\times$ benchmark cells). Thus, where tested byte-exactly, the ranking rather than the forward path explains the gap (§5). This matched-rule result is distinct from the query-conditioned FastV comparison: FastV remains stronger on several scene-text/object-centric cells, so RBM is a query-blind, OCR-preserving robust default, not a universal winner (§4.4).

Contributions. First, we isolate a selection-level mechanism on Qwen3-VL: rank reshuffling, text-stroke demotion, and ranking-swap control identify the kept set as the lever (§5). Second, we establish a workload-conditioned stage law across three models and two merger designs under the shared L2 family and iso-budget control (§4). Third, we derive minimal RBM, map its trade-off against FastV, and show that six prespecified extensions do not improve their stronger constituents (§4.4, §6).

2 Related Work

Token pruning and merging in VLMs. Most methods score visual tokens and drop or merge the least important before or within the early LLM layers. **FastV** prunes at the second LLM layer from the attention visual tokens receive [6]; it is our primary same-model, query-conditioned strong baseline — the method that is stronger on several question-conditioned non-OCR cells but weaker in the tested OCR regime (§4.4). **PyramidDrop** drops tokens progressively across LLM layers [39]; **FasterVLM** accelerates attention-based pruning [48]; **SparseVLM** explores query-conditioned sparsification [49]; **PruMerge** merges by attention-derived importance [29]; **FitPrune** fits lightweight importance predictors [42]. Other training-free selectors use visual [CLS] attention, graph propagation, or multi-cue suppression [14, 25, 34], while unified recent systems combine importance with diversity or text-guided merging [1, 10, 43]. These methods prune or merge at different encoder-output or LLM-side hook points. Among the methods reviewed here, we found no matched comparison of selection immediately before versus after a model’s native spatial

Kept visual units (blue) per selector — two OCR cases




Figure 1: Three selectors retain different visual evidence. Two audited OCRBench cases compare RBM, Post-L2, and FastV-k3 at the same $\kappa = 25\%$ visual-token budget on Qwen3-VL-8B. Blue native-grid cells are the visual units retained by each method; the underlying input remains visible for direct spatial comparison. Vector checks and crosses report answer correctness against the printed ground truth. In both examples, RBM preserves the answer-bearing text and answers correctly, while Post-L2 and FastV-k3 do not. These cases illustrate the selection behavior; aggregate effects and uncertainty are reported in Figure 2.

merger under iso-model, iso-token, and a shared query-blind L2-family rule.

Token-merging lineage (ToMe $\rightarrow$ AdaptMerge). Token merging originates with ToMe [4] (bipartite soft matching, post-encoder, query-blind); **AdaptMerge** [13] adapts it to large multimodal models and reports that language-guided merging closes most of the OCR gap plain ToMe incurs — corroborating evidence for our lossy-merger thesis, cited as support, not as a competitor. This lineage adds a *second*, learned merging stage; we hold the native merger fixed and ask only whether selection should precede or follow it, so it is methodologically distinct from our stage-axis test. Learned visual projectors such as TokenPacker condense tokens while preserving fine local cues [21]; unlike our inference-time test, they alter the projector architecture and require model adaptation.

VisionZip and Qwen-family compression. VisionZip selects “dominant” tokens by CLS-to-patch attention and merges the remainder into “contextual” tokens [41]. In the official Qwen2.5-VL implementation we inspected, selection operates on embeddings after the PatchMerger, i.e. on the post-merger side of our axis. Its authors’ own Qwen2.5-VL numbers are a model/stage-mismatched

anchor only: OCRBench 81.5 $\rightarrow$ 70.5 at 50% retention while general tasks hold, with the README caveat that gains are “less striking” because the merger already compresses — the same effect seen from the post-merger side; we claim no head-to-head victory over it (§4.4). Qwen-specific methods **GlimpsePrune** [44] and **VS-can** [45] provide additional Qwen-family compression comparisons without the matched stage-only control used here. The concurrent **AnchorPrune** study applies Qwen2.5-VL selection after its model-specific spatial processing and token merging [28], again placing it on the post-merger side rather than matching both sides of the boundary. Recent layerwise methods instead score token evolution: TransPrune uses token-transition variation [19], while V2Drop exploits variation across adjacent layers [5]. They are complementary to the controlled native-merger stage axis studied here.

The 2026 frontier. Recent work broadens token reduction beyond scalar attention ranks through hierarchical selection [33]. Set coverage, reconstruction error, and functional sensitivity motivate joint or projector-aware selection [9, 15, 17, 26, 40]; learned or strongly conditioned compressors instead model semantic labels, information flow, or head importance [31, 32, 37, 51]. Other work

makes token lifetime explicitly layer-dependent [7, 22, 36], or preserves discarded content through object-centric, optimal-transport, or hybrid aggregation [18, 20, 46]. Document- and OCR-specific studies further show that generic accuracy can hide lost textual evidence [8, 23]. These advances reinforce the need to localize the intervention stage; unlike these studies, our experiment directly matches the two sides of a frozen model’s native spatial merger.

The design point we isolate. We do not claim a topological vacuum. QuietPrune-style in-ViT pruning [11] is a different, earlier, trained stage; PixelPrune removes redundant pixel patches before the ViT [35]; FastV/PyramidDrop select inside the LLM; VisionZip selects on the opposite (post-merger) side. The combination we isolate is narrower: on a merger-equipped VLM, leaving the native merger untouched, scoring the merger-*input* 2×2 units by a query-blind saliency norm, keeping the top- κ , and handing only the survivors to the unmodified native merger — together with Qwen3-VL mechanism evidence (M1–M3) that is not evaluated in our cited comparisons and explains why this stage can dominate on text-dense workloads.

Efficiency evaluation of VLMs. Reported speedups depend heavily on engine and protocol. We adopt the Imms-eval family of *official* scorers for accuracy [47]; EffiVLM-Bench independently emphasizes cross-backbone, multi-metric acceleration evaluation [38]. We confine all throughput claims to a single engine (§7).

3 Method — Rank-Before-Merge (RBM)

3.1 Models, mergers, and the stage axis

We study three merger-equipped VLMs with two merger designs. Qwen3-VL-8B and Qwen2.5-VL-7B use a learned PatchMerger [2, 3]; InternVL3-8B uses pixel-shuffle downsampling [50]. All are served in bf16, eager attention, vLLM 0.19 (V1) [16], single A40 (46 GB), greedy decoding. The shared pipeline is: image $\rightarrow$ ViT patch tokens $\rightarrow$ native merger (2×2 groups $\rightarrow$ one token) $\rightarrow$ language model. The Qwen merger is *not* an averaging pool: the four patch features of each 2×2 group are concatenated and projected by a small learned MLP (LayerNorm $\rightarrow$ fc1 $\rightarrow$ GELU $\rightarrow$ fc2) to one unit vector; the lossy step is the four-to-one nonlinear projection, so a kept unit’s representation is a nonlinear function of its four patches. We verified this against the served implementation on both Qwen generations.

The stage axis (the experimental variable). The two arms share model, token count, decoding, and a query-blind L2 magnitude rule, while evaluating that rule on the native feature state available on each side of the merger:

- **Pre-merger (RBM).** For each 2×2 unit, average the L2 norms of its four merger-input patch vectors, keep the top- κN , and pass only the survivors through the unmodified native merger (and, for Qwen3-VL, each deepstack merger). The merger operates exactly as in the uncompressed model, on a subset of units.
- **Post stage (contrast).** Either score the N merged units on merger-*output* features and keep the top- κN (the stage used by published compression methods for merger-equipped VLMs, including VisionZip’s Qwen path), or — the stronger, query-conditioned contrast — prune inside the LLM by

layer-2 attention (the FastV stage). The main matrix (Table 1) uses the post-merger-L2 contrast from the same query-blind L2 family; the layer-2 FastV contrast is used only in Table 2.

Budget definition. Retention κ is defined over *merge units* relative to each image’s **own** full unit count, so two methods at the same κ feed the language model the same number of visual tokens per image; we verify equality of the mean post-merger token count per benchmark (**iso-token** control), so any accuracy difference is attributable to *which* units survive and the representational state on which selection is made, not to token count. We report $\kappa \in \{0.25, 0.125\}$ (25% / 12.5% retention); per-benchmark token counts are printed with each table.

3.2 The L2 selector as a control variable

We deliberately use the simplest **text-agnostic, query-blind** L2-magnitude family. For unit u with four merger-input patch vectors $h_{u,1:4}$ and native merger M , the implemented scores are

$$s_{\text{pre}}(u) = \frac{1}{4} \sum_{j=1}^4 \|h_{u,j}\|_2, \quad s_{\text{post}}(u) = \|M(h_{u,1:4})\|_2.$$

Thus the ranking functional is matched at the level of feature magnitude, while its representation is the experimental stage: four input vectors before the nonlinear merger versus its output vector after it. For Qwen3-VL, the headline operational pre arm takes h from the first deepstack tap (ViT layer 8), whereas the matched-depth pre-final control takes h from the final ViT block at the main-merger boundary (§4.2). We freeze plain RBM and add no deployed variant. §5 additionally reports a mechanism-side probe: a *variance-weighted* variant of the same query-blind scorer, $s_{\text{freq}}(u) = \alpha z(\|f_u\|_2) + \beta z(\text{Var}_u)$ (both terms per-image standardized, so the ratio $\alpha:\beta$ is a genuine trade-off knob rather than an arbitrary scale), improves text-dense accuracy under aggressive compression (+1.2 pp TextVQA at 25% retention, +0.5 pp at 50%) yet is not a general upgrade: it regresses -1.6 to -2.2 pp on DocVQA, OCR-Bench, and GQA at 25%. We therefore treat it as corroborating evidence for §5.2’s text-stroke hypothesis rather than as a decoration of the frozen selector (§6). §5.4 reports that the tested effect survives replacing L2 with a second scorer family (global-centroid attention) on Qwen3-VL, indicating the effect is not an L2 artifact. Unlike the merging lineage of §2, RBM introduces no second learned merger and no query input: it manipulates only the *set* of units the native merger sees — which is exactly what makes the pre-vs-post contrast a clean stage experiment rather than a scorer comparison.

3.3 Rank-Before-Merge

Given an image with N merge units, compute $s(u)$ on merger-input features for every unit, retain the top $k = \kappa N$, and invoke the native merger(s) on the retained units only. Unit identity is shared across the main and deepstack streams (deepstack mergers process intermediate-depth features of the *same* units), so pruning a unit removes it from every stream. No attention weights are required and no parameters are trained.

Unit equivalence (the structural invariant). Because the mask is at *unit* granularity and the merger is per-unit, a fixed

Algorithm 1 Rank-Before-Merge (RBM), one image**Require:** image x ; retention κ ; query-blind scorer $s = \|\cdot\|_2$ **Ensure:** compressed visual token sequence z

```

1:  $P \leftarrow$  ViT patch tokens of  $x$  {16 px footprint}
2:  $U \leftarrow$  group  $P$  into  $N$  native  $2 \times 2$  units {32 px footprint}
3: for each unit  $u \in U$  do
4:    $\text{score}[u] \leftarrow \frac{1}{4} \sum_{j=1}^4 \|h_{u,j}\|_2$  {mean patch L2}
5: end for
6:  $k \leftarrow \max(1, \text{round}(\kappa N))$ ;  $K \leftarrow \text{topk}(\text{score}, k)$ 
7:  $z_{\text{main}} \leftarrow \text{NativeMerger}(\{u : u \in K\})$  {merge survivors only}
8:  $z_{\text{ds}} \leftarrow \text{DeepstackMergers}(\{u : u \in K\})$  {Qwen3-VL only}
9:  $z \leftarrow \text{NativeVisualInterface}(z_{\text{main}}, z_{\text{ds}})$ 
10: return  $z$  {feed standard LLM}

```

kept set supplies the same four patches to the same native per-unit operation. Selection cannot prevent intra-unit combination; it changes which units reach the merger. The resulting outputs are byte-identical in the verified Qwen3-VL path; other families are limited to the evidence stated in §8. Pre-merger selection keeps units on their raw (merger-input) saliency; post-stage selection keeps them on a saliency the merger has already rewritten (§5). This invariant is what makes the ranking-swap control of §5.3 decisive.

Post-merger selection (the contrast): replace lines 3–5 of Algorithm 1 by $U' \leftarrow \text{NativeMerger}(U)$; $\text{score}'[u'] \leftarrow \|f_{\text{out}}(u')\|_2$ on merger-*output* features; $K' \leftarrow \text{top-}k$ by score' ; keep(K'). Every-thing downstream is identical.

Configuration details. Per-family tap points (Qwen3-VL deepstack[0] at ViT layer 8, Qwen2.5-VL main merger input, InternVL3 pixel-shuffle input) and the Qwen2.5-VL position-cursor fix are detailed in Supplementary §S8, including the cross-generation configuration diagnostic.

4 Main Experiments

4.1 Setup

Benchmarks and metrics. We use TextVQA VQA accuracy [30], DocVQA ANLS [27], the official OCR-Bench score [24], and GQA normalized exact match [12]. Official split sizes are 5000, 5349, 1000, and 12578, respectively. Scorers are verbatim ports of the lmms-eval family [47], with ground-truth self-test passing 200/200.

Engines and fairness protocol. RBM and the none / post cells run in vLLM; the FastV baseline runs in an independent Hugging-Face transformers eager harness (it has no vLLM path). Accuracy is comparable across engines (the harness bit-degrades to native at zero drop, 8/8 per-sample identical; HF-vs-vLLM none cells agree 16/16); **throughput is not** — all efficiency numbers are confined to vLLM (§7). Fairness is enforced by the same keep ratio relative to each image’s own token count and by reporting mean post-merger token counts.

Significance. Each pre-vs-post Δ pairs the same samples under greedy decoding. The prespecified primary family comprises the nine model $\times$ text-dense-benchmark pre-vs-post contrasts at 25% retention in Table 1. We use 20,000-resample paired-bootstrap confidence intervals on per-item metric differences and two-sided paired sign-flip permutation tests; Holm correction across these

nine tests gives $p_{\text{Holm}} \leq 4.5 \times 10^{-4}$ for every contrast. McNemar is reported only as a binary-outcome cross-check for OCRBench containment and GQA exact match. Other budgets, GQA, mechanism probes, FastV comparisons, and extension gates are secondary or exploratory and are not pooled into the primary family. The metric-level tests preserve partial credit in VQA accuracy and ANLS rather than binarizing those metrics.

4.2 Main result across three models — official full splits (Table 1)

The stage law holds on both generations. Every text-dense Δ_{25} is large and significant: RBM leads post-merger selection by +38.4 pp (TextVQA), +24.3 pp (DocVQA), +363 pts (OCR-Bench; 547 vs 184, 3.0 $\times$) on Qwen3-VL, and by +26.1 pp, +11.0 pp, +293 pts (476 vs 183, 2.6 $\times$) on Qwen2.5-VL; all paired-bootstrap 95% CIs exclude zero. The gap is robust to deeper compression: at $\kappa = 0.125$ the post stage degrades much faster (Qwen3-VL OCR-Bench 184 $\rightarrow$ 53 vs RBM 547 $\rightarrow$ 350, widening the pre/post ratio from 3.0 $\times$ to 6.6 $\times$), and all six text-dense $\Delta_{12.5}$ remain significant under paired metric-level inference.

A pure-stage confound control (pre-final). The original pre arm on Qwen3-VL ranks at the deepstack[0]-input tap (layer 8, far upstream) and slices all four mergers, while the post arm ranks at the main-merger *output* — so the pre-vs-post contrast differs in *both* stage and feature space. To isolate the stage alone we run a **pre-final** control that ranks L2 at the main-merger’s *own input* (final ViT-block features, the same feature space the post arm ranks *after* the merger) and selects the top- κ merge-units from the visual output, leaving the deepstack mergers untouched ($n = 200$, Qwen3-VL, greedy, $\kappa = 0.25$, L2, iso-config with Table 1; iso-token verified, 200/200 per benchmark). The only difference between pre-final and post is then *before versus after* the lossy 2×2 main merger — a pure stage variable.

Pre-final still beats post by +30.0 pp on TextVQA and +12.5 pp on DocVQA — the same direction and order as the original pre-vs-post gap (+38.3/+26.5 pp at $n = 200$). The full control table and paired uncertainty are in Supplementary S7b. The text-dense pre $\gg$ post claim is therefore a **true stage effect** (ranking before the lossy merger preserves text), *not* a deepstack[0]-input versus main-merger-output feature-space artifact; the pure-stage control is valid (iso-token verified).

Object-centric GQA — no detected pure-stage difference. On GQA, pre-final *equals* post (0.450 = 0.450, 0.0 pp, paired mean 0.000 ± 0.032): this experiment detects no difference between ranking before and after the merger on object-QA, consistent with the mechanism (the merger does not preferentially damage object features; M2 $\rho = 0.04$ on GQA, §5). The original pre-vs-post contrast on Qwen GQA (−2.8 pp Qwen3-VL / −2.6 pp Qwen2.5-VL at full $n = 12578$; primary McNemar $z = -8.0/-5.7$; independent-binomial $z = -4.5/-4.1$; 83.9% / 83.8% per-question agreement; the official exact-match rescore gives a concordant −8.1/−7.1) is therefore *not* a true stage reversal: it is an artifact of the deepstack[0]-input (layer-8) ranking being poorly suited to object-QA, dissolved once the ranking taps the main-merger input (pre-final, 0.0 pp). Its magnitude is an order of magnitude below any text-dense Δ (1/10–1/14 of them), and **no text-dense benchmark shows any crossover**

Table 1: Official full-split metrics ($n = 5000/5349/1000/12578$), greedy, across Qwen3-VL-8B, Qwen2.5-VL-7B, and InternVL3-8B. RBM = pre-merger L2 (ours); post = post-merger L2 at the same token budget. $\Delta_{25} = \text{pre} - \text{post}$ at $\kappa = 0.25$. † marks the only significant post-stage leads (GQA); these use the deepstack[0]-input (layer-8) pre ranking and are reframed as a layer-8 ranking artifact by the pure-stage control in §4.2 (pre-final gives 0.0 pp on GQA). InternVL3 uses a different merger and family-specific imaging configuration; compare within-model deltas, not cross-model absolute values. Full protocol and skip accounting are in the supplementary material.

Model	Benchmark (metric)	none	RBM @25%	post @25%	Δ_{25}	RBM @12.5%	post @12.5%
Qwen3-VL-8B	TextVQA (VQA-acc)	0.844	0.605	0.222	+38.4 pp	0.472	0.132
	DocVQA (ANLS)	0.956	0.481	0.238	+24.3 pp	0.352	0.103
	OCR-Bench (/1000)	760	547	184	+363 pts	350	53
	GQA (exact match)	0.616	0.449	0.477	−2.8 pp †	—	—
Qwen2.5-VL-7B	TextVQA (VQA-acc)	0.862	0.702	0.442	+26.1 pp	0.597	0.319
	DocVQA (ANLS)	0.949	0.636	0.526	+11.0 pp	0.455	0.245
	OCR-Bench (/1000)	817	476	183	+293 pts	335	67
	GQA (exact match)	0.604	0.559	0.585	−2.6 pp †	—	—
InternVL3-8B	TextVQA (VQA-acc)	0.8338	0.7890	0.4148	+37.4 pp	0.7230	0.3064
	DocVQA (ANLS)	0.9221	0.7284	0.3820	+34.6 pp	0.5054	0.2451
	OCR-Bench (/1000)	852	753	321	+432 pts	—	—
	GQA (exact match)	0.6293	0.5993	0.6031	−0.4 pp (n.s.)	—	—

Table notes. (a) Qwen3-VL none DocVQA was re-run at max-model-len 49,152 and completed with 0/5349 skips (ANLS 0.956); the pre/post cells it anchors ran at the same native-resolution configuration and are directly comparable to each other. (b) OCR-Bench none cells skip 18 (Qwen3-VL) / 24 (Qwen2.5-VL) of 1000 long-OCR images on context overrun, scored 0 (conservative; compressed cells skip ≤ 5), so the uncompressed anchor is understated and the relative gaps are, if anything, conservative. (c) GQA was evaluated at $\kappa = 0.25$ only on the full split. (d) Iso-token control: at each model×benchmark, pre and post have identical mean post-merger token counts, with one exception — the Qwen2.5-VL OCR-Bench @25% RBM cell used a 4 M-pixel cap (mean tokens 229.6 vs post 282.0), so its +293-pt gap is conservative. (e) Primary inference uses paired-bootstrap 95% CIs and paired sign-flip permutation tests on the official per-item metric. All Qwen text-dense intervals exclude zero; Figure 2 shows the intervals. McNemar is additional for binary OCRBench/GQA outcomes. (f) InternVL3 uses pixel-shuffle downsampling and a 4 M-pixel DocVQA cap; its pre/post prompt token counts are identical within every cell. Its GQA −0.4 pp difference is indistinguishable: the paired 95% CI [−0.93, +0.17] crosses 0 (permutation $p = 0.18$), so we claim no GQA stage lead on InternVL3; no preregistered equivalence bound exists, so we do not call it a statistical tie.

at the native-resolution full splits. We therefore claim **no detected pure-stage difference** on this object-centric sample; this is not an equivalence test. All other sections refer here for the GQA statistics.

A no-merger negative control (LLaVA-1.5). LLaVA-1.5-7B has a per-token MLP projector (Linear–GELU–Linear, no 2x2 spatial averaging). Running the same pre-vs-post L2 protocol (pre-MLP vs post-MLP ranking, keep 25%, $n=200$, greedy) yields *no* significant stage effect: TextVQA $\Delta = -2.17$ pp (McNemar $p=0.125$; none 0.452, pre 0.087, post 0.108), GQA $\Delta = -0.5$ pp ($p=1.0$; none 0.575, pre 0.375, post 0.380) — versus the merger-equipped Qwen3-VL’s +38.4 pp ($p \approx 5 \times 10^{-5}$) on TextVQA. In this no-merger control the trend is slightly reversed (post $\geq$ pre, n.s.). Together with the merger-equipped results, this is consistent with the spatial merger being a key condition for the large text-dense pre $>$ post gap. A single negative-control architecture cannot establish necessity or an if-and-only-if law.

4.3 Replication on a third family — InternVL3-8B (Table 1)

To test whether the Qwen direction replicates beyond that family, we run RBM and the post stage on **InternVL3-8B**, whose merger is pixel-shuffle downsampling — a different design from Qwen’s PatchMerger. The post arm here is **post-merger L2 top-k** — the same query-blind L2-magnitude family as the pre arm, at an identical token budget (ptid equal across arms), so the InternVL3 block of

Table 1 isolates the before/after merger stage under iso-budget control at 25% retention. The query-conditioned in-layer FastV contrast is a separate campaign, reported in Table 2.

Configuration boundary. InternVL3 DocVQA uses a 4 M-pixel cap, while the Qwen same-scope baseline cells of §4.4 use 600k; cross-model absolute values are not comparable. At 25% retention RBM costs −19.4 pp vs none on InternVL3 DocVQA and −99/1000 on OCR-Bench, versus only −4.5 pp on TextVQA and −3.0 pp on GQA. The claim is therefore the large within-model pre $>$ post margin, not losslessness on every text-dense workload.

At 12.5% retention ($r = 0.875$), TextVQA is 0.7230 pre vs post 0.3064 (**+41.7 pp**); DocVQA pre 0.5054 vs post 0.2451 (**+26.0 pp**) — the gap *widens* under deeper compression, as on Qwen.

Cross-family directional replication. InternVL3’s pre–post text-dense gaps (+37.4/+34.6 pp / +432 pts) are the same order and direction as Qwen3-VL’s (+38.4/+24.3/+363 pts) on all three text-dense benchmarks, and GQA stays within the red line — indistinguishable here (−0.4 pp) versus Qwen3-VL’s apparent post edge (+2.8 pp, a deepstack[0]-input layer-8 ranking artifact dissolved by the pure-stage control, §4.2), with no text-dense crossover at either family’s native imaging configuration. These accuracy-level replications show that the shared L2-magnitude rule can favor the pre-merger side on a third model family with a different merger. They support — but do not by themselves establish — the Qwen3-VL selection-level mechanism as a cross-family explanation.

The ranking-swap comparison on InternVL3. The M3-style control (§5) — run the post-merger forward path but select units

with the pre-merger ranking — closely matches pre-merger accuracy on the three tested InternVL3 cells: swap vs pre is +0.00 pp (TextVQA), +0.11 pp (DocVQA), and −0.50 pp (GQA), with kept-set Jaccard(pre,swap)= 1.0 and per-sample answer-identity 0.98–1.0 ($n = 200$, iso-token, no swap fallback). This supports a ranking-level explanation on a second merger architecture (pixel-shuffle, not just Qwen’s PatchMerger), but the nonzero residuals mean that the byte-exact causal attribution remains specific to Qwen3-VL.

4.4 A workload-conditioned regime map — RBM vs the query-conditioned strong baseline FastV (Table 2)

FastV is the query-conditioned strong baseline in this comparison — it selects inside the LLM after cross-attention has mixed in the question — and we give it every advantage: we compare RBM against FastV at its **best layer** K ($K = 3$; a K -sensitivity probe found $K = 3$ best), reported at identical n /scope with paired skips and official rescoring, so the comparison cannot be attributed to a handicapped K or a mismatched sample set. We read the result as a **workload-conditioned regime map** — where query-conditioned selection wins, where it fails, and why — not as a single-winner race. At its canonical progressive schedule PyramidDrop is near-lossless but occupies a different budget point (equivalent mean retention $0.625 \approx 2.1\times$ ours); forced to an honest iso-25% schedule it collapses (TextVQA 0.073 / GQA 0.305 / OCR 0.005 on Qwen3-VL), whereas RBM holds (0.60 / 0.42 / 0.58). There is therefore no universal winner.

Workload-conditioned regime map. FastV leads several question-conditioned cells. FastV-k3 leads RBM on all three Qwen3-VL non-OCR cells (+17.2/+8.9/+16.2 pp) and on Qwen2.5-VL TextVQA (+7.8 pp). This is the query-dependent headroom made reachable: *which* region matters for *this question* becomes visible only once cross-attention has seen the question, exactly where FastV operates and exactly where the prespecified query-level gate of §6 could not reach before the merger. **RBM leads the tested OCR regime.** RBM wins OCR-Bench on both families (+16.0 pp Qwen3-VL, +8.5 pp Qwen2.5-VL); the Qwen2.5-VL $n = 200$ GQA/DocVQA cells favor RBM by +4.5/+2.1 pp but their paired 95% CIs cross 0 ($[-10.5, +1.5]$ / $[-10.1, +5.8]$), so they are exploratory/inconclusive at $n = 200$, not statistically robust RBM edges. The observed OCR split persists at FastV’s best tested K ; within this comparison, later query conditioning does not recover the OCR accuracy retained by RBM. **The map is the two failure modes of §5.4.** Documents = selection misdirection — the merger’s reshuffle misleads selection made after it away from text, so the OCR regime is lost by the selectors that tap the stream post-merger; scene text = feature degradation under a budget every 25%-keep method pays, and object scenes = near-null reshuffle — together the low-OCR regime where the query-dependent headroom lives. We report FastV’s wins honestly: FastV is a strong query-conditioned baseline that wins its regime; RBM is the robust default (never collapses, uniquely holds dense OCR), not a method that leads query-conditioned pruning everywhere.

Table 2: Workload-conditioned regime map — same-scope paired comparison of FastV at best $K = 3$ (the query-conditioned strong baseline) vs RBM (pre, the query-blind robust default) @25% retention, both Qwen models, official rescoring. Cells are paired within each row (identical n , identical attempted samples; skip parity in note ii). Margin = winner’s absolute lead; the winner column is the regime read: FastV’s wins mark its reference-grounded home regime, RBM’s wins mark the OCR regime (§5.4).

Benchmark (scope; pixel cap)	n	FastV-k3	RBM @25%	Δ (F/R)
Qwen3-VL-8B				
TextVQA (full; default)	5000	0.7771	0.605	+17.2 F
GQA (full; default)	12578	0.5376	0.449	+8.9 F
DocVQA (dev; 600k cap)	200	0.5863	0.4239	+16.2 F
OCR-Bench (dev; native res.)	200	0.415	0.575	+16.0 R
Qwen2.5-VL-7B				
TextVQA (dev; default)	200	0.7467	0.6683	+7.8 F
GQA (dev; default)	200	0.475	0.520	+4.5 R [‡]
DocVQA (dev; 600k cap)	200	0.4852	0.5062	+2.1 R [‡]
OCR-Bench (dev; native res.)	200	0.285	0.370	+8.5 R

Table notes. F = FastV; R = RBM; Δ is the winner’s absolute lead in percentage points.

(i) Scope and pixels: n is printed per row; DocVQA cells run at a 600k pixel cap on both arms (iso-pixel within this table); OCR-Bench runs at native resolution, which is what triggers the skips (note ii); TextVQA/GQA use each model’s default imaging configuration. (ii) OCR-Bench skips: native-resolution giant OCR images overrun the context on $\approx 10\%$ of the $n = 200$ set (19/200 Qwen3-VL, 20/200 Qwen2.5-VL); the identical hard samples are skipped under both arms (same paired batch, mode-independent) and scores are over attempted samples — conservative and symmetric. All other cells skip ≤ 5 (Qwen3-VL TextVQA FastV, 5/5000) or 0. (iii) Harness: every $n = 200$ cell on both sides comes from the same HF transformers eager-attention harness with family-correct native positions; the Qwen3-VL TextVQA/GQA RBM references are the vLLM full-split cells of Table 1, for which engine parity is validated (§4.1; HF-pre vs vLLM-pre @25% further agree 14/16 with kept-set Jaccard 0.986, §6d). An early Qwen2.5-VL $\times$ pre position-harness degeneration (vllm-mimic layout) was detected and replaced by the native layout before any reported number was produced; no reported cell comes from the degenerate path (on Qwen3-VL, native vs mimic differ, 0.575 vs 0.525, same ordering). (iv) Budget: FastV-k3 ≈ 214 vs RBM ≈ 213 mean visual tokens/request (Qwen3-VL TextVQA, 25% retention; the $n = 200$ probe measured 208) — approximately matched, stated per cell; this is “each method at its strong configuration,” not a byte-iso-budget cell. (v) [‡]The paired 95% CI crosses 0 at $n = 200$ (GQA $[-10.5, +1.5]$, DocVQA $[-10.1, +5.8]$); these two cells are exploratory/inconclusive, not statistically robust RBM edges.

5 Mechanism — Why the Merger Is Lossy

We give the pre-vs-post gap a mechanistic account in three parts (M1–M3), measured on **Qwen3-VL-8B** (M1/M2 on a deterministic seed-0 $n = 64$ sample per benchmark; M3 on $n = 200$ with the same invocation as the headline cells), then state its architectural scope.

5.1 M1 — the merger reshuffles unit ranks

Per image we correlate the *pre* ranking (merger-input unit L2) with the *post* ranking (merged-token L2) over all units. The two rankings are barely related, and the reshuffle tracks text density: Spearman $\rho = 0.137$ (DocVQA, text-dense document) < 0.332 (TextVQA, text-dense scene) < 0.360 (GQA, object-centric), with the decision-relevant top-25% overlap rising in the same order (Jaccard@25% = 0.180 $< 0.243 < 0.278$; full values in Supplementary S1). The merger’s nonlinear projection thus reorders unit saliency almost from scratch ($\rho = 0.14$ – 0.36 , nowhere near 1); both the full-ranking




Figure 2: The text-dense stage effect replicates across models and merger designs. Points are paired pre-post metric differences at 25% retention; whiskers are paired-bootstrap 95% confidence intervals on the official full splits. OCRBench’s /1000 score is divided by 10 for a common percentage-point axis. Qwen uses PatchMerger; InternVL3 uses pixel shuffle. GQA bounds the result: Qwen’s original layer-8 pre-ranking favors post, while InternVL3’s interval crosses zero. Values correspond to Table 1.

correlation and the decision-relevant top-25% overlap rise from the most text-dense benchmark to the object-centric one, so the reshuffle is largest on DocVQA and smallest on GQA — yet substantial on every benchmark.

5.2 M2 — the demoted units are text-stroke units

With $\text{rank_shift} = \text{post_rank} - \text{pre_rank}$ ($+$ $\Rightarrow$ demoted) and per-unit Sobel edge energy as a text-stroke proxy, the merger preferentially *demotes* high-edge units at $\kappa = 25\%$: $\rho(\text{rank_shift}, \text{edge}) = +0.439$ on DocVQA, $+0.155$ on TextVQA, $+0.036$ on GQA — $\sim 12\times$ stronger on documents than on object scenes. The units post drops that pre keeps are the highest-edge of all: on DocVQA, mean Sobel **0.641 vs 0.124** for the reverse group, with **92% above the per-image median (vs 35%)** — post-merger selection is **systematically anti-text on documents**. (The *direction*, we hypothesize, is set by the merger’s training on natural content, which attenuates high-frequency stroke energy; this is a testable hypothesis, not a measured result.) The full M2 table (per-benchmark Sobel means and median fractions) is in Supplementary S2.

5.3 M3 — a ranking-swap causal control, and the decisive kept-set identity

By unit equivalence (§3.3), holding the *forward path* fixed at post and swapping in the *pre ranking* must reproduce pre accuracy if the gap is purely a ranking effect. It does:

Swapping in the pre ranking recovers full pre-merger accuracy (exactly on DocVQA; within 2/200 greedy-decode answers on TextVQA, ϵ -level kernel numerics confirmed by a 200/200 byte-identical independent rerun; full values are in Supplementary S9) and erases the collapse ($+38.3$ pp TextVQA, $+26.5$ pp DocVQA over post). The swap path agrees with post on only 40/200 (TextVQA)

and 63/200 (DocVQA) answers — without the pre ranking, the same forward path produces the collapse.

The decisive hardening — kept-set identity on both architectures. Emitting the kept unit indices from both paths and comparing directly: $\text{Jaccard}(\text{swap-kept}, \text{pre-top-}\kappa) = 1.000$ in all four cells (Qwen3-VL and Qwen2.5-VL $\times$ DocVQA and TextVQA; $n = 32$, chance level $\kappa/N = 0.25$). The two paths select *exactly the same units*. The small swap-vs-pre answer residual on Qwen2.5-VL ($+6.7$ pp DocVQA at $n = 32$) traces to an implementation confound of the swap control — window-attention reverse indices (three recovery sites on Qwen2.5-VL, zero on Qwen3-VL) map identical index sets to physical units in different orderings — while batch-dependent merging is structurally excluded (the merger is per-token LN+MLP). **Verdict: on Qwen3-VL, the tested pre-post gap is attributable to the ranking alone at selection level** (byte-identical); Qwen2.5-VL corroborates kept-set identity but retains an ordering-confound residual. The merged representations of kept units carry no stage-dependent loss on Qwen3-VL (byte-exact); on Qwen2.5-VL the kept *set* is identical but the ordering confound is not eliminated, so we scope the no-loss statement to the byte-exact family.

5.4 Failure modes and scope boundaries

DocVQA exhibits selection misdirection: post ranking moves away from text and the swap fully restores accuracy. TextVQA combines rank corruption with a budget effect, while GQA has near-null text directionality ($\rho = 0.04$). The stage effect is therefore workload-conditional. It retains sign and magnitude under a global-centroid attention scorer on Qwen3-VL, but that proxy degrades on Qwen2.5-VL; L2 remains the paper’s fixed selector. The ranking law replicates on Qwen2.5-VL and at the accuracy level on InternVL3 (Table 1),

whereas the causal decomposition is claimed only for byte-exact Qwen3-VL. A low 600k pixel cap reverses shallow Qwen2.5-VL DocVQA; the native-resolution full split remains the headline configuration (§8).

Retention curve. The stages are nearly indistinguishable at 75% retention; their gap then widens monotonically through 25% on both Qwen models (full values in Supplementary S3, $n = 200$).

Merger mechanism. Figure 3 connects the stage effect to the native merger. Across four workloads, pre- and post-merger rankings share only moderate rank correlation and top-quartile overlap. Units rich in text-stroke edge energy are preferentially demoted after merging. Finally, a Qwen3-VL ranking-swap control keeps the selected unit set fixed (Jaccard = 1.000) while restoring the pre-merger ordering, separating ranking damage from a change in token identity.

6 Negative Results — Tested Extensions Do Not Compose

We prespecified six query-aware, regime-aware, or composite extensions and report all six as failures on Qwen3-VL-8B. **(a) Image-level routing:** pre is at least as accurate as post on 84–97% of images per benchmark; the best router scores below always-pre (0.484 vs 0.494) and far below the per-sample oracle (0.576). On pooled $N = 774$, only ~27% of the oracle gap is workload-level, leaving ~73% sample- and query-dependent. **(b) Merger-aware masking:** its selected $t = 0.5$ hybrid gains 2.3 pp on TextVQA but loses 8.0 pp on OCR-Bench and gains nothing on GQA; no candidate passes the no-OCR-regression plus GQA-gain gate. **(c) Query-embedding blending:** on a disjoint development slice every $\lambda > 0$ underperforms plain L2 (−1.7/−5.2/−3.3 pp for $\lambda = 0.3/0.5/0.7$), so the prespecified rule selects $\lambda = 0$. **(d) RBM→FastV cascade:** for $n = 200 \times 4$ benchmark samples and two budgets, it misses the 0.5-pp Pareto margin in all eight cells, typically by 8–18 pp, and never significantly beats both arms. Complete gates, paired counts, and harness controls are in Supplementary S4b/S9.3. **(e) A variance-weighted (frequency-aware) selector:** the mechanism-side probe of §3.2 ($\alpha=1$, $\beta=0.6$, selected in two grid rounds on TextVQA-500 at 25% retention) improves the text-dense cell (+1.2 pp TextVQA@25%, +0.5 pp at 50%) but regresses −1.6 to −2.2 pp on DocVQA@25%, OCR-Bench@25%, and GQA@25%; no candidate passes a no-regression gate across the four benchmarks, so we freeze the L2 selector and treat the improvement as hypothesis support only. **(f) An adaptive PRE/POST stage router:** a query-blind workload detector on merger-input features (high-frequency tail ratio + L2-entropy) routes each image to the stage predicted to be stronger. Tuned on TextVQA it reduces to always-PRE — statistically indistinguishable from plain RBM on every benchmark (TextVQA/DocVQA/OCR-Bench regression = 0.0 pp, but GQA gain = +0.0 pp, below the required +1 pp); raising the threshold to route scene images to POST collapses TextVQA by a further 7 pp. The GQA and TextVQA workload signatures overlap at the unit level, so no single global threshold reconciles the two acceptance criteria.

Exploratory learned boundary scorer. A held-out MLP probe successfully distills post-merger importance from pre-merger norm, edge, variance, position, and neighborhood features, but transfers the wrong task bias: at 25% retention, TextVQA falls from 0.595 to

0.435 (−16.0 pp) while GQA rises from 0.515 to 0.550 (+3.5 pp). It is therefore not an RBM upgrade; as an intervention, its failure is consistent with M2’s text-stroke demotion (Supplementary S4c.3).

These results bound, rather than close, the tested design space: image-only signals cannot access the query-dependent oracle headroom, cascading discards evidence needed by the second selector, and distilling a post-merger teacher reproduces its text-dense failure mode. We therefore freeze plain RBM as the robust text-dense/OCR default and select a FastV-class method for scene-text/object-centric traffic; within the tested configurations, routing, cascading, and learned post-order imitation offer no consistent benefit.

7 Efficiency — Supporting Evidence

Efficiency is supporting evidence, measured only in vLLM 0.19 V1 offline on one A40. **Stage-neutral:** across 75/50/25% retention, pre vs post throughput is 4.61/4.52, 5.39/5.54, and 6.91/6.89 req/s; five-repeat measurements remain within 3% with no consistent direction. **Compression scaling:** relative to no compression, throughput rises 12/31/68% and 25%-retention latency falls from 0.36 to 0.23 s. InternVL3 independently yields 1.8–2.5× throughput at 25% retention, again with pre $\approx$ post at identical token counts. The flat ≈ 42 GB allocator reservation is not a memory measurement; continuous-batching efficiency remains future work.

8 Discussion and Limitations

Across three tested models, pre-merger selection better preserves text-dense accuracy; Qwen3-VL diagnostics indicate that its native merger rewrites saliency in a text-hostile direction. RBM is therefore a robust default, not a universal winner. It retains dense OCR better than the compared FastV settings (+16.0/+8.5 pp on the two Qwen models), while FastV leads on several scene-text/object-centric cells (Table 2).

Architecture and mechanism scope. The byte-exact M3 causal attribution is established on Qwen3-VL; Qwen2.5-VL retains an ordering confound, and InternVL3 provides accuracy-level and bounded ranking-swap support rather than a byte-exact causal decomposition. Models without a native merger are out of scope. M1/M2 use a Sobel proxy on $n = 64$ samples per benchmark and M3 uses $n = 200$; the GQA null has no equivalence margin. Holm correction covers the nine prespecified 25%-retention text-dense stage contrasts; secondary budgets, GQA, mechanism analyses, baseline regime comparisons, and extension gates remain uncorrected and should be read as secondary or exploratory.

Evaluation scope. FastV uses HF eager whereas the other methods use vLLM; anchors support accuracy comparison, but throughput is vLLM-only. Some baseline cells use $n = 200$, as marked. Metric-level bootstrap intervals and paired permutation tests are primary; McNemar covers additional binary outcomes.

Configuration scope. DocVQA pixel caps differ (4 M for InternVL3, 600k for Qwen baseline cells), so cross-model absolute scores are incomparable. Low-resolution Qwen2.5-VL reverses the shallow-stage direction; the headline is the native-resolution full split. RBM also remains lossy against no compression on InternVL3 DocVQA/OCR (−19.4 pp/−99 points), despite beating post.




Figure 3: What the native merger changes. (a) Pre/post L2 rankings are only moderately preserved, measured by Spearman correlation and top-25% kept-set Jaccard across DocVQA ($n = 64$), OCRBench ($n = 199$), TextVQA ($n = 64$), and GQA ($n = 64$). (b) Rank shifts are associated with text-stroke edge energy: units with stronger high-frequency structure are demoted more after merging. (c) In the Qwen3-VL ranking-swap control, post-merger features use the pre-merger order while the kept set remains identical (Jaccard = 1.000), isolating ordering from token identity on $n = 200$ TextVQA and DocVQA samples.

9 Conclusion

We isolate visual-token selection stage under iso-model, iso-budget, and a shared query-blind L2-magnitude family. Across the tested text-dense cells of three models from two architecture families, pre-merger ranking is consistently stronger; the measured gaps generally grow under deeper compression. A ranking-swap control attributes the Qwen3-VL gap to merger-induced rank changes. The resulting training-free method, RBM, is a robust OCR-oriented default rather than a universal winner: query-conditioned FastV remains stronger on several scene-text/object-centric settings. RBM adds no stage-specific serving cost, and six prespecified Qwen3-VL extensions fail to improve over their stronger constituents.

References

- [1] Changwoo Baek, Jouwon Song, Sohyeon Kim, and Kyeongbo Kong. 2026. AgilePruner: An Empirical Study of Attention and Diversity for Adaptive Visual Token Pruning in Large Vision-Language Models. In *International Conference on Learning Representations (ICLR)*. arXiv:2603.01236
- [2] Shuai Bai, Yuxuan Cai, Ruizhe Chen, Keqin Chen, Xionghui Chen, Zesen Cheng, Lianghao Deng, Wei Ding, Chang Gao, Chunjiang Ge, et al. 2025. Qwen3-VL Technical Report. arXiv:2511.21631 [cs.CV]
- [3] Shuai Bai, Keqin Chen, Xuejing Liu, Jialin Wang, Wenbin Ge, Sibao Song, Kai Dang, Peng Wang, Yuanzhi Zhu, Junyang Lin, et al. 2025. Qwen2.5-VL Technical Report. arXiv:2502.13923 [cs.CV]
- [4] Daniel Bolya, Cheng-Yang Fu, Xiaoliang Dai, Peizhao Zhang, Christoph Feichtenhofer, and Judy Hoffman. 2023. Token Merging: Your ViT But Faster. In *International Conference on Learning Representations (ICLR)*. arXiv:2210.09461 [cs.CV]
- [5] Junjie Chen, Xuyang Liu, Zichen Wen, Yiyu Wang, Siteng Huang, and Honggang Chen. 2026. Variation-Aware Vision Token Dropping for Faster Large Vision-Language Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 3489–3499. arXiv:2509.01552
- [6] Liang Chen, Haozhe Zhao, Tianyu Liu, Shuai Bai, Junyang Lin, Chang Zhou, and Baobao Chang. 2024. An Image is Worth 1/2 Tokens After Layer 2: Plug-and-Play Inference Acceleration for Large Vision-Language Models. In *European Conference on Computer Vision (ECCV)*. arXiv:2403.06764 [cs.CV]
- [7] Yongru Chen, Kai Zhang, Zeliang Zong, Yuchen Lu, Wenming Tan, Ye Ren, and Jilin Hu. 2026. One Layer’s Trash is Another Layer’s Treasure: Adaptive Layer-wise Visual Token Selection in LVLs. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*.
- [8] Joonmyung Choi, Sanghyeok Lee, Jongha Kim, Sehyung Kim, Dohwan Ko, Ji-hyung Kil, and Hyunwoo J. Kim. 2026. DocPrune: Efficient Document Question Answering via Background, Question, and Comprehension-aware Token Pruning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 3543–3552.
- [9] Chenxi Du, Yongheng Deng, Jiani Liu, Yujia Zhang, Xi Chen, and Ju Ren. 2026. CoIn: Coverage and Informativeness-Guided Token Reduction for Efficient Large Multimodal Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 10492–10501.
- [10] Zhengyao Fang, Pengyuan Lyu, Chengquan Zhang, Guangming Lu, Jun Yu, and Wenjie Pei. 2026. Prune Redundancy, Preserve Essence: Vision Token Compression in VLMs via Synergistic Importance-Diversity. In *International Conference on Learning Representations (ICLR)*. arXiv:2603.09480
- [11] Tianxiao Gao, Shanwei Zhao, Shuo Fang, Shuai Zhu, and Chenguang Ma. 2026. QuietPrune: Query-Guided Early Token Pruning for Vision-Language Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 3553–3562.
- [12] Drew A. Hudson and Christopher D. Manning. 2019. GQA: A New Dataset for Real-World Visual Reasoning and Compositional Question Answering. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. arXiv:1902.09506 [cs.CV]
- [13] Zahidul Islam and Mrigank Rochan. 2025. AdaptMerge: Inference Time Adaptive Visual and Language-Guided Token Merging for Efficient Large Multimodal Models. In *Findings of the Association for Computational Linguistics: EMNLP 2025*. doi:10.18653/v1/2025.findings-emnlp.387
- [14] Yutao Jiang, Qiong Wu, Wenhao Lin, Wei Yu, and Yiyi Zhou. 2025. What Kind of Visual Tokens Do We Need? Training-free Visual Token Pruning for Multi-modal Large Language Models from the Perspective of Graph. arXiv:2501.02268 [cs.CV]
- [15] Youngeun Kim, Youjia Zhang, Huiling Liu, Aechon Jung, Sunwoo Lee, and Sungeun Hong. 2026. ZOO-Prune: Training-Free Token Pruning via Zeroth-Order Gradient Estimation in Vision-Language Models. In *Proceedings of the*

- IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 39572–39582. arXiv:2509.24837
- [16] Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. 2023. Efficient Memory Management for Large Language Model Serving with PagedAttention. In *Proceedings of the 29th Symposium on Operating Systems Principles (SOSP)*. arXiv:2309.06180 [cs.OS]
 - [17] Jaeyeon Lee, Shunjie Wen, and Dong-Wan Choi. 2026. Moving Beyond Diversity: Visual Token Pruning as Subspace Reconstruction for Efficient VLMs. arXiv:2606.18681 [cs.CV]
 - [18] Jingyu Lei, Gaoang Wang, and Der-Hong Lee. 2026. CORE: Compact Object-centric REpresentations as a New Paradigm for Token Merging in LLMs. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*.
 - [19] Ao Li, Yuxiang Duan, Jinghui Zhang, Congbo Ma, Yutong Xie, Gustavo Carneiro, Mohammad Yaqub, and Hu Wang. 2026. TransPrune: Token Transition Pruning for Efficient Large Vision-Language Model. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 39529–39538. arXiv:2507.20630
 - [20] Jinlong Li, Liyuan Jiang, Haonan Zhang, and Nicu Sebe. 2026. Token Reduction via Local and Global Contexts Optimization for Efficient Video Large Language Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. arXiv:2603.01400
 - [21] Wentong Li, Yuchuan Yuan, Jian Liu, Dongqi Tang, Song Wang, Jie Qin, Jianke Zhu, and Lei Zhang. 2024. TokenPacker: Efficient Visual Projector for Multimodal LLM. arXiv:2407.02392 [cs.CV]
 - [22] Xu Li, Yi Zheng, Mengyang Zhao, Yuxuan Liang, Zhe Liu, Rui Zhu, Xiaolei Chen, Wei Zhou, Baoquan Zhao, and Juncen Guo. 2026. CRISP: Pre-LLM Yet Text-Driven Visual Token Pruning for Efficient LLM Inference. arXiv:2607.16326 [cs.CV]
 - [23] Feixiang Liu, Qiang Qiu, Hao Zhang, and Xinyue Wang. 2026. Beyond Accuracy: Auditing Spatial Provenance in Visual Token Pruning for OCR-Critical MLLM Inference. arXiv:2608.00077 [cs.CV]
 - [24] Yuliang Liu, Zhang Li, Mingxin Huang, Biao Yang, Wenwen Yu, Chunyuan Li, Xucheng Yin, Cheng-Lin Liu, Lianwen Jin, and Xiang Bai. 2024. OCRBench: On the Hidden Mystery of OCR in Large Multimodal Models. *Science China Information Sciences* 67, 22 (2024), 220102. arXiv:2305.07895 [cs.CV] doi:10.1007/s11432-024-4235-6
 - [25] Bozhi Luan, Wengang Zhou, Hao Feng, Zhe Wang, Xiaosong Li, and Houqiang Li. 2025. Multi-Cue Adaptive Visual Token Pruning for Large Vision-Language Models. arXiv:2503.08019 [cs.CV]
 - [26] Qiankun Ma, Ziyao Zhang, Haoqi Wang, Zhen Song, Jie Chen, and Hairong Zheng. 2026. ApET: Approximation-Error Guided Token Compression for Efficient VLMs. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*.
 - [27] Minesh Mathew, Dimosthenis Karatzas, R. Manmatha, and C. V. Jawahar. 2021. DocVQA: A Dataset for VQA on Document Images. In *IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*. 2200–2209. arXiv:2007.00398 [cs.CV]
 - [28] Kyuan Oh and Bumsoo Kim. 2026. AnchorPrune: Relevance-Anchored Contextual Expansion for Visual Token Pruning. In *European Conference on Computer Vision (ECCV)*. arXiv:2607.07033
 - [29] Yuzhang Shang, Mu Cai, Bingxin Xu, Yong Jae Lee, and Yan Yan. 2025. LLaVA-Prune: Adaptive Token Reduction for Efficient Large Multimodal Models. In *IEEE/CVF International Conference on Computer Vision (ICCV)*. arXiv:2403.15388 [cs.CV] doi:10.1109/ICCV51701.2025.02122
 - [30] Amanpreet Singh, Vivek Natarajan, Meet Shah, Yu Jiang, Xinlei Chen, Dhruv Batra, Devi Parikh, and Marcus Rohrbach. 2019. Towards VQA Models That Can Read. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. arXiv:1904.08920 [cs.CV]
 - [31] Jiafei Song, Fengwei Zhou, Jin Qu, Wenjin Jason Li, Tong Wu, Gengjian Xue, Zhikang Zhao, Daomin Wei, Yichao Lu, and Bailin Na. 2026. EvoComp: Learning Visual Token Compression for Multimodal Large Language Models via Semantic-Guided Evolutionary Labeling. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*.
 - [32] Guohao Sun, Yufei Wang, Sizhuo Ma, Yuege Xie, Yuting Cheng, Zhiqiang Tao, and Jian Wang. 2026. IF-Prune: Information-Flow Guided Token Pruning for Efficient Vision-Language Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 3522–3531.
 - [33] Zixun Sun, Yubo Dong, Hehe Fan, and Yi Yang. 2026. Hi-Lo Prune: Look at What You'll Lose before Pruning with Hierarchical Token Selection. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 31941–31951.
 - [34] Ao Wang, Fengyuan Sun, Hui Chen, Zijia Lin, Jungong Han, and Guiguang Ding. 2024. [CLS] Token Tells Everything Needed for Training-free Efficient MLLMs. arXiv:2412.05819 [cs.CV]
 - [35] Nan Wang, Zhiwei Jin, Chen Chen, and Haonan Lu. 2026. PixelPrune: Pixel-Level Adaptive Visual Token Reduction via Predictive Coding. arXiv:2604.00886 [cs.CV]
 - [36] Yahong Wang, Juncheng Wu, Zhangkai Ni, Longzhen Yang, Yihang Liu, Chengmei Yang, Ying Wen, Lianghua He, Xianfeng Tang, Hui Liu, and Yuyin Zhou. 2026. When Token Pruning is Worse than Random: Understanding Visual Token Information in VLLMs. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 31910–31919. arXiv:2512.07580
 - [37] Yi Wang, Haoqi Zhang, Qihan Huang, Anda Cao, Gongfan Fang, Wei Wang, Xuan Jin, Jie Song, Mingli Song, and Xinchao Wang. 2026. Rethinking Token Reduction for Large Vision-Language Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 24727–24737. arXiv:2603.21701
 - [38] Zekun Wang, Minghua Ma, Zexin Wang, Rongchuan Mu, Liping Shan, Ming Liu, and Bing Qin. 2025. EffiVLM-BENCH: A Comprehensive Benchmark for Evaluating Training-Free Acceleration in Large Vision-Language Models. In *Annual Meeting of the Association for Computational Linguistics (ACL)*. arXiv:2506.00479
 - [39] Long Xing, Qidong Huang, Xiaoyi Dong, Jiajie Lu, Pan Zhang, Yuhang Zang, Yuhang Cao, Conghui He, Jiaqi Wang, Feng Wu, and Dahua Lin. 2025. PyramidDrop: Accelerating Your Large Vision-Language Models via Pyramid Visual Redundancy Reduction. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. arXiv:2410.17247 [cs.CV]
 - [40] Tong Xu, Hailong Shi, and Xingyu Gao. 2026. SCoRe: Saliency-Coverage Reduction for Vision Token Pruning in Vision-Language Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*.
 - [41] Senqiao Yang, Yukang Chen, Zhuotao Tian, Chengyao Wang, Jingyao Li, Bei Yu, and Jiaya Jia. 2025. VisionZip: Longer is Better but Not Necessary in Vision Language Models. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. arXiv:2412.04467 [cs.CV]
 - [42] Weihao Ye, Qiong Wu, Wenhao Lin, and Yiyi Zhou. 2025. Fit and Prune: Fast and Training-free Visual Token Pruning for Multi-modal Large Language Models. In *Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)*. arXiv:2409.10197 [cs.CV]
 - [43] Hanxun Yu, Wentong Li, Xuan Qu, Song Wang, Junbo Chen, and Jianke Zhu. 2026. VisionTrim: Unified Vision Token Compression for Training-Free MLLM Acceleration. In *International Conference on Learning Representations (ICLR)*. arXiv:2601.22674
 - [44] Quan-Sheng Zeng, Yunheng Li, Qilong Wang, Peng-Tao Jiang, Zuxuan Wu, Ming-Ming Cheng, and Qibin Hou. 2026. A Glimpse to Compress: Dynamic Visual Token Pruning for Large Vision-Language Models. *IEEE Transactions on Circuits and Systems for Video Technology* (2026), 1–1. arXiv:2508.01548 [cs.CV] doi:10.1109/TCSVT.2026.3702147
 - [45] Ce Zhang, Kaixin Ma, Tianqing Fang, Wenhao Yu, Hongming Zhang, Zhisong Zhang, Haitao Mi, and Dong Yu. 2025. VScan: Rethinking Visual Token Reduction for Efficient Large Vision-Language Models. arXiv:2505.22654 [cs.CV]
 - [46] Jusheng Zhang, Xiaoyang Guo, Kaitong Cai, Qinhan Lv, Yijia Fan, Wenhao Chai, Jian Wang, and Keze Wang. 2026. Hybrid Token Compression for Vision-Language Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 31889–31899. arXiv:2512.08240
 - [47] Kaichen Zhang, Bo Li, Peiyuan Zhang, Fanyi Pu, Joshua Adrian Cahyono, Kairui Hu, Shuai Liu, Yuanhan Zhang, Jingkang Yang, Chunyuan Li, and Ziwei Liu. 2025. LLMs-Eval: Reality Check on the Evaluation of Large Multimodal Models. In *Findings of the Association for Computational Linguistics: NAACL 2025*. 881–916. arXiv:2407.12772 [cs.CV]
 - [48] Qizhe Zhang, Aosheng Cheng, Ming Lu, Renrui Zhang, Zhiyong Zhuo, Jiajun Cao, Shaobo Guo, Qi She, and Shanghang Zhang. 2024. Beyond Text-Visual Attention: Exploiting Visual Cues for Effective Token Pruning in VLMs. arXiv:2412.01818 [cs.CV]
 - [49] Yuan Zhang, Chun-Kai Fan, Junpeng Ma, Wenzhao Zheng, Tao Huang, Kuan Cheng, Denis Gudovskiy, Tomoyuki Okuno, Yohei Nakata, Kurt Keutzer, and Shanghang Zhang. 2025. SparseVLM: Visual Token Sparsification for Efficient Vision-Language Model Inference. In *International Conference on Machine Learning (ICML) (Proceedings of Machine Learning Research, Vol. 267)*. PMLR. arXiv:2410.04417 [cs.CV]
 - [50] Jinguo Zhu, Weiyun Wang, Zhe Chen, Zhao Yang, Shenglong Ye, Lixin Gu, Hao Tian, Yuchen Duan, Weijie Su, Jie Shao, Zhangwei Gao, Erfei Cui, Xuehui Wang, Yue Cao, Yangzhou Liu, Xingguang Wei, Hongjie Zhang, Haomin Wang, Weiye Xu, Hao Li, Jiahao Wang, Nianchen Deng, Songze Li, Yanan He, Tan Jiang, Jiapeng Luo, Yi Wang, Conghui He, Botian Shi, Xingcheng Zhang, Wenqi Shao, Junjun He, Yingting Xiong, Wenwen Qu, Peng Sun, Penglong Jiao, Han Lv, Lijun Yu, Kaipeng Zhang, Huipeng Deng, Jiaye Ge, Kai Chen, Limin Wang, Min Dou, Lewei Lu, Xizhou Zhu, Tong Lu, Dahua Lin, Yu Qiao, Jifeng Dai, and Wenhao Wang. 2025. InternVL3: Exploring Advanced Training and Test-Time Recipes for Open-Source Multimodal Models. arXiv:2504.10479 [cs.CV] doi:10.48550/arXiv.2504.10479
 - [51] Qihui Zhu, Tao Zhang, Yuchen Wang, Shuangwu Chen, Xiaobin Tan, Jian Yang, Yang Liu, and Yinfei Pan. 2026. HAWK: Head Importance-Aware Visual Token Pruning in Multimodal Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*.

Supplementary Material

Rank Before You Merge: A Stage Law for Training-Free Visual Token Compression in Merger-Equipped Vision-Language Models

Every number is indexed by an audit identifier in S9. Tables S1–S3/S7 support § 5 of the body; S4–S6 support § 4 (S4b cascade counts support § 6); S8 supports § 3; S9 gives the per-cell artifact index, with cell, sample count, official metric, token count, skip accounting, and manifest identifier in S9.1–S9.4; S10 is a qualitative same-image pre/post selection comparison rendered from measured Qwen3-VL-8B merger features (four panels each, no paired question).

S1. M1 — merger ranking reshuffle, full table (Qwen3-VL-8B, seed-0, $n = 64$ /benchmark)

Per image, Spearman correlation of the pre ranking (merger-input unit L2) with the post ranking (merged-token L2) over all units, plus top-25% kept-set overlap. Evidence record: audit index entry for the M1 mechanism campaign. The Kendall- τ column is taken from the same audited M1 record.

Benchmark	Spearman ρ	Jaccard@25%	Kendall τ
DocVQA (document)	0.137	0.180	0.094
TextVQA (scene)	0.332	0.243	0.227
GQA (object)	0.360	0.278	0.249

Reading: the reshuffle is largest on the most text-dense benchmark and smallest (but still substantial) on object-centric GQA — $\rho = 0.137 < 0.332 < 0.360$; Jaccard = $0.180 < 0.243 < 0.278$. For two independent size- κN subsets, the large- N chance reference is $\kappa/(2 - \kappa) = 1/7 \approx 0.143$ at $\kappa = 25\%$. All three overlaps are above chance but low in absolute terms, with DocVQA closest to chance. Per-image distributions are in Supplementary Figure 5 (top).

S2. M2 — demotion directionality vs text-stroke energy, full table (same sample as S1)

$\text{rank_shift} = \text{post_rank} - \text{pre_rank}$ ($+$ $\Rightarrow$ demoted); per-unit Sobel edge energy as the text-stroke proxy. Evidence record: anonymized artifact manifest entry for the M2 mechanism campaign. The per-benchmark Sobel means and median fractions are taken from the same audited M2 record.

Benchmark	ρ	Sobel drop	Sobel reverse	demoted > median
DocVQA	+0.439	0.641	0.124	92% (vs 35%)
TextVQA	+0.155	0.281	0.186	72% (vs 53%)
GQA	+0.036	0.304	0.271	58% (vs 53%)

Reading: the merger preferentially demotes high-edge (text-stroke) units $\sim 12\times$ more strongly on documents than on object scenes; on DocVQA the units post drops that pre keeps are the highest-edge of all (92% above the per-image median vs 35% for the reverse group) — post-merger selection is systematically anti-text on documents. Directionality hypothesis (merger trained on natural content attenuates high-frequency stroke energy): testable hypothesis, not a measured result. Second text-stroke proxy (Laplacian variance): future work (single-proxy disclosure).

S3. Retention curve — full number strings ($n = 200$, both Qwen models)

pre-post Δ at 75% / 50% / 25% retention (the body § 5.4 prints endpoints only). Evidence record: audit index entry for the retention-curve ablation.

Model	Benchmark	Δ @75%	Δ @50%	Δ @25%
Qwen3-VL-8B	TextVQA	+8.7 pp	+29.0 pp	+38.3 pp
Qwen3-VL-8B	DocVQA	−1.3 pp	+5.9 pp	+26.5 pp
Qwen2.5-VL-7B	TextVQA	+7.0 pp	+15.3 pp	+32.0 pp
Qwen2.5-VL-7B	DocVQA	−2.1 pp	−1.4 pp	+18.8 pp

Reading: the two stages are indistinguishable under shallow compression (DocVQA post even leads 1–2 pp within noise at 75%); the pre-merger lead emerges and widens monotonically with depth. Same shallow regime as the Qwen2.5-VL 600k-cap reversal (body § 5.4 boundary; § 8 item 3). The Qwen3-VL curves are visualized in Supplementary Figure 5 (bottom).




Figure 4: Representative token-survival maps at 25% retention. TextVQA (top) and DocVQA (bottom) use the same query-blind L2 selector before and after Qwen3-VL’s native merger. In each row, the first two panels show the pre- and post-selected units; the third overlays pre-kept (green), post-kept (red dashed), and pre-only/post-dropped units (orange). These are audited single-image mechanism examples, not aggregate accuracy evidence.




Figure 5: Quantitative mechanism and depth diagnostics. Top: per-image pre/post rank agreement on the deterministic seed-0 M1 sample ($n = 64$ per benchmark); boxes and points show the distributions, with mean±standard deviation annotated. Bottom: normalized accuracy as compression deepens on Qwen3-VL-8B ($n = 200$ per point). Pre-merger selection preserves substantially more text-dense accuracy, while ChartQA is primarily budget-limited and GQA does not exhibit the text-dense stage effect. Error bars are descriptive rate-approximation standard errors; primary inference uses paired metric-level bootstrap and sign-flip tests.

S4. Paired McNemar cross-checks — per cell (full splits, both Qwen models)

Binary-outcome cross-check = paired McNemar z on per-sample correctness, $z = (b - c) / \sqrt{b + c}$, with b/c the pre-only / post-only correct counts; positive z = pre-merger lead, negative = post-merger lead (GQA only). Cross-check = independent-binomial z ($SE \sqrt{se_{pre}^2 + se_{post}^2}$); “agree %” = per-question agreement of the recorded correctness indicators. Evidence record: audit index entry for the paired-test table; paired by question id, restricted to ids answered in both arms; the 200/200 ground-truth self-test is retained in the audit record.

† the only significant post-stage lead. These GQA cells use the deepstack[0]-input (layer-8) pre ranking; a pure-stage control that ranks at the main-merger input (pre-final) gives 0.0 pp on GQA (body § 4.2, S7b), so the apparent post>pre is a layer-8 ranking artifact, not a true stage reversal. GQA official word-normalized exact-match rescore gives a concordant McNemar $z = -8.1 / -7.1$ (Qwen3-VL / Qwen2.5-VL). Smallest text-dense $\Delta_{25} |z| = 14.6$ (Qwen2.5-VL OCR-Bench); smallest text-dense $\Delta_{12.5} |z| = 15.0$ (Qwen2.5-VL OCR-Bench); smallest $|z|$ anywhere = 5.7 (Qwen2.5-VL GQA). On OCR-Bench @25% the independent-binomial cross-check (18.2 / 14.7) marginally exceeds McNemar (16.7 / 14.6) because the binary correctness indicator coarsens the /1000 item score (body footnote e); McNemar remains a cross-check and is $\geq$ -consistent on every other cell.

S4b. Cascade two-stage gate — paired McNemar counts (Qwen3-VL-8B, $n = 200$, official rescore)

Cascade gate (body § 6): single budget, Stage 1 = pre-merger RBM keeping fraction X of merger-input units, Stage 2 = FastV at LLM layer 2 dropping fraction Y of the survivors; total keep = $X \cdot (1 - Y)$. Three arms (pre-alone / fastv-alone / cascade) are compared pairwise on identical images (n_{common} matched; mean post-merger token count identical across all three arms in every cell, so the comparison is iso-budget). cas_only / pre_only / fst_only = per-sample counts correct in only the named arm of the pair. Verdict = NO-GO (cascade fails the Pareto rule on all 8 budget×benchmark cells). Evidence record: audit index entry for the cascade gate and paired McNemar counts.

Reading: cascade sits strictly between the two single methods and is dominated by the better arm on every cell — pre-alone on OCR-Bench ($z = -5.15 / -5.01$ vs pre at 25% / 12.5%; the FastV stage destroys exactly the raw-patch information Stage 1 was built to protect) and FastV-alone on TextVQA / GQA ($z = -2.34 / -2.43$ vs fst at 25%). The single positive cas-vs-pre z (DocVQA 25%, +2.07) is paired with a loss to FastV (−1.50), so the “strictly beats both arms” condition holds nowhere. Two training-free selectors composed in series do not complement (body § 6).

S4c. Selector-scorer and stage-router gates (Qwen3-VL-8B, official rescore)

S4c.1 Frequency-aware selector (body § 6e): the mechanism-side probe of body § 3, $s_{\text{freq}}(u) = \alpha z(\|f_u\|_2) + \beta z(\text{Var}_u)$, tuned in two grid rounds on TextVQA₅₀₀ at 25% retention ($\alpha=1, \beta=0.6$; official VQA-acc). The variance term helps only the most text-dense cell and only under aggressive compression; everywhere else it is neutral or negative, so no no-regression-across-four-benchmarks gate passes. Evidence record: anonymized audit index entry for the frequency-aware selector.

S4c.2 Adaptive stage router (body § 6f): query-blind per-image workload detector on merger-input features — high-frequency tail ratio $r_{\text{HF}} = \Pr[\text{Var}_u > \bar{\text{Var}} + \sigma_{\text{Var}}]$ and 20-bin L2-entropy $H(\|f_u\|_2)$ — routing each image to PRE (RBM, rank-before-merge) or POST (post-merger L2), both iso-budget. κ is the kept fraction; the winner of a $r_{\text{HF}} \times \tau_{\text{ent}}$ grid on TextVQA₅₀₀ ($r_{\text{HF}}=0.08, \tau_{\text{ent}}=2.0$) routes every image to PRE, reproducing plain RBM bit-for-bit: text benchmarks show 0.00 pp regression but GQA gain is also +0.00 pp (below the required +1 pp). The $r_{\text{HF}}=0.13$ threshold does route some scene images to POST but collapses TextVQA (51/320 flipped, −7 pp) while flipping only 22/243 GQA images (0.428 vs RBM 0.438). GQA and TextVQA r_{HF} distributions overlap (means 0.127 vs 0.142), so no single global threshold reconciles the two acceptance criteria. Evidence record: audit index entry for the adaptive stage-router gate and workload-signal analysis.

S4c.3 Exploratory learned boundary scorer (body § 6): we test whether a light-weight learned selector can improve on plain L2 without moving selection past the merger. A small MLP receives only pre-merger boundary features available at inference time: unit L2 norm, Sobel edge energy, local feature variance, normalized position, and neighborhood statistics. Its teacher is the post-merger importance ordering captured on separate training images; evaluation uses disjoint held-out benchmark samples and the official task scorers. Both learned and L2 arms retain exactly 25% of merger-input units and otherwise share the same Qwen3-VL-8B inference configuration.

The asymmetric task result is diagnostic rather than a method improvement. M2 independently shows that the post-merger ordering demotes text-stroke-rich units; distillation transfers this preference into a pre-merger scorer and causes a 16.0-pp TextVQA loss. The positive GQA result shows that the learned head is capable of

changing task behavior, ruling out a trivial “model did not learn” explanation. We therefore retain plain L2 and do not use the learned scorer in any headline table.

S5. Descriptive rate-scale errors + skip accounting (main full-split Qwen blocks)

For binary exact-match/containment outcomes, $\sigma = \sqrt{p(1-p)/n}$ is the Bernoulli standard error. For VQA accuracy and ANLS, which retain partial credit, applying the same expression to the printed mean is only a descriptive rate-scale approximation, not an inferential error bar. All claims use the paired metric-level bootstrap and sign-flip tests in the main paper. Evidence record: audit index entry for the Qwen full-split table and skip accounting.

- Qwen3-VL none DocVQA: re-run at max-model-len 49,152, native resolution, **0/5349 skips** (ANLS 0.956).
- OCR-Bench none cells skip **18/1000 (Qwen3-VL) / 24/1000 (Qwen2.5-VL)** long-OCR images on context overrun, scored 0 (conservative); compressed cells skip ≤ 5 . The uncompressed anchor is therefore understated and the relative gaps are conservative.
- Qwen2.5-VL OCR-Bench @25% RBM: 4 M-pixel cap (mean tokens 229.6 vs post 282.0) $\rightarrow$ its +293-pt gap is conservative (the one iso-token exception in the Qwen block of Table 1, note d).

Per-cell descriptive σ values (full splits, computed from the printed official metric p and full-split n):

These values provide scale and skip-accounting context only; they are not used to test partial-credit metric differences. The full-split GQA post lead is an order of magnitude smaller than the text-dense gaps and is reframed by the pre-final control (S7b) as a layer-8 ranking artifact rather than a true stage reversal. ^a Qwen3-VL none-DocVQA native-resolution anchor: full-split run partly skipped on context overrun (re-run at max-model-len 49,152 gives ANLS 0.956, 0/5349 skips); σ not printed for the partial native cell. ^c GQA evaluated at $\kappa = 0.25$ only on the full split (no full-split $\kappa = 0.125$ cells; body footnote c).

- Qwen3-VL OCR-Bench none = 760/1000 over 982 attempted (18 skips scored 0); compressed @25%: 547 (pre) / 184 (post), skip ≤ 5 — ratio 3.0×, widening to 6.6× at 12.5% (350 vs 53).

S6. $n = 200$ cross-split consistency (dev-subset vs full-split direction)

The $n = 200$ dev scope used by the regime-map and mechanism probes is direction-consistent with the full splits. Evidence record: audit index entry for the cross-split consistency campaign. $\Delta_{25} = \text{pre} - \text{post}$ at $\kappa = 0.25$.

Reading: all four text-dense cells agree in sign (pre-merger lead) across the two scopes; magnitudes are within run noise (the $n = 200$ subset SE is $\approx 8\times$ larger, so the subset is retained only as a direction check, never as the headline). The full Qwen3-VL $n = 200$ cross-method view (RBM vs FastV vs Pyramid) and the compact Qwen2.5-VL $n = 200$ cross-generation matrix appear in the main paper’s regime table and the archived audit index, respectively.

S7. Selector invariance (second scorer family, Qwen3-VL)

The stage law under a second, non-L2 scorer (global-centroid attention). Evidence record: audit index entry for the selector-invariance campaign.

Reading: invariant in sign and magnitude on Qwen3-VL — the effect is not an L2 artifact. On Qwen2.5-VL the L2 sign is invariant but the attention proxy degrades (proxy-family specificity, not a counter-example); L2 remains the paper’s selector, frozen as plain RBM (§ 3.2).

S7b. Pure-stage confound control — pre-final (Qwen3-VL-8B, $n = 200$)

The body § 4.2 pre-final control isolates the *stage* from the feature-space confound of the original pre (deepstack[0]-input, layer 8) vs post (main-merger output) contrast. pre-final ranks L2 at the main-merger’s own input (final ViT-block features) and selects the top- κ merge-units from the visual output; the deepstack mergers run untouched. The only difference from post is then before vs after the lossy 2×2 main merger — a pure stage variable. Iso-token verified (prompt token ids match 200/200 per benchmark for pre-final vs post). Evidence record: anonymized audit index entry for the pre-final control and its per-sample audit.

Reading: on text-dense benchmarks pre-final still leads post by a large margin (the text-dense stage effect is a true stage effect, not a deepstack[0]-input vs main-merger-output feature-space artifact). On GQA pre-final equals post (0.0 pp): there is no stage effect on object-QA, and the original pre-vs-post apparent post>pre (−2.8 pp Qwen3-VL, −2.6 pp Qwen2.5-VL; S4) was an artifact of the deepstack[0]-input layer-8 ranking being poorly suited to object-QA, not a true stage reversal. Consistent with the mechanism: the merger does not preferentially damage object features (M2

Model	Benchmark	κ	Δ (official)	McNemar z (check)	b (pre-only)	c (post-only)	agree %	indep-binom z
Qwen3-VL-8B	TextVQA	.25	+38.4 pp	+43.0	2349	183	49.4	42.3
Qwen3-VL-8B	DocVQA	.25	+24.3 pp	+34.7	1601	150	67.3	27.1
Qwen3-VL-8B	OCR-Bench	.25	+363 pts	+16.7	422	56	52.2	18.2
Qwen3-VL-8B	GQA	.25	-2.8 pp $\dagger$	-8.0	834	1196	83.9	-4.5
Qwen3-VL-8B	TextVQA	.125	+34.0 pp	+40.8	2101	161	54.8	39.9
Qwen3-VL-8B	DocVQA	.125	+24.9 pp	+33.1	1437	127	70.8	32.1
Qwen3-VL-8B	OCR-Bench	.125	+297 pts	+16.0	323	25	65.2	17.8
Qwen2.5-VL-7B	TextVQA	.25	+26.1 pp	+30.8	1682	309	60.2	27.3
Qwen2.5-VL-7B	DocVQA	.25	+11.0 pp	+15.6	1406	693	60.8	11.6
Qwen2.5-VL-7B	OCR-Bench	.25	+293 pts	+14.6	347	54	59.9	14.7
Qwen2.5-VL-7B	GQA	.25	-2.6 pp $\dagger$	-5.7	892	1150	83.8	-4.1
Qwen2.5-VL-7B	TextVQA	.125	+27.8 pp	+33.1	1836	305	57.2	29.1
Qwen2.5-VL-7B	DocVQA	.125	+21.0 pp	+25.6	1325	294	69.7	23.3
Qwen2.5-VL-7B	OCR-Bench	.125	+268 pts	+15.0	294	26	68.0	15.9

total keep κ	bench	n	pre	fst	cas	z (cas vs pre)	cas_only	pre_only	z (cas vs fst)	cas_only	fst_only
25%	TextVQA	200	0.5967	0.6800	0.5950	-0.152	21	22	-2.335	18	35
25%	DocVQA	200	0.4239	0.5183	0.4779	+2.066	38	22	-1.500	26	38
25%	OCR-Bench	181	0.5801	0.1713	0.3978	-5.154	4	37	+6.112	43	2
25%	GQA	200	0.4150	0.4900	0.4050	-0.333	17	19	-2.429	16	33
12.5%	TextVQA	200	0.5000	0.4450	0.4050	-2.540	21	41	-0.896	27	34
12.5%	DocVQA	200	0.2980	0.2779	0.2070	-2.012	31	49	-2.064	16	30
12.5%	OCR-Bench	181	0.3591	0.0884	0.1713	-5.013	6	40	+3.441	17	2
12.5%	GQA	200	0.3950	0.4550	0.3500	-1.286	20	29	-3.130	12	33

bench@ κ	freq($\alpha=1, \beta=0.6$)	RBM (L2)	Post-L2	Δ freq-RBM
TextVQA @25%	0.6173	0.6053	0.2073	+1.20
TextVQA @50%	0.6780	0.6733	0.3873	+0.47
DocVQA @25%	0.4903	0.5063	0.2317	-1.60
DocVQA @50%	0.6271	0.6212	0.5487	+0.59
OCR-Bench @25%	0.7160	0.7380	0.1700	-2.20
OCR-Bench @50%	0.8220	0.8220	0.4980	0.00
GQA @25%	0.4220	0.4380	0.4540	-1.60
GQA @50%	0.4860	0.4840	0.5000	+0.20

Table 3: Frequency-aware selector (official metrics, $n = 500$; TextVQA=1.2 pp gain at $\kappa=25\%$, baseline-equal elsewhere, strong regression on the other three benchmarks at 25%).

Benchmark	Learned POST teacher	Plain L2	Δ learned-L2
TextVQA	0.435	0.595	-16.0 pp
GQA	0.550	0.515	+3.5 pp

Table 4: Held-out learned-scorer evaluation at 25% retention ($n = 200$ per benchmark; official metrics). The scorer can imitate a post-merger teacher, but the resulting ordering is strongly anti-task on TextVQA and therefore fails the cross-workload no-regression gate.

$\rho = 0.04$ on GQA, S2). pre-final tested on Qwen3-VL only (the deepstack confound is Qwen3-specific; Qwen2.5-VL shares the architecture but was not re-tested).

S8. Merger tap points and M-RoPE configuration details (supports §3)

- **Qwen3-VL-8B**: the first deepstack merger has index 0 and the visual indexes are [8, 16, 24], so unit scores use **ViT block-8 outputs** (intermediate depth); the single resulting mask is shared by all mergers, including the main

merger that consumes the final-block output. The tap-point and call-order verification is retained in the audit index.

- **Qwen2.5-VL-7B**: no deepstack; scores are computed on the **main merger input (final ViT block output)**. mask_computed_at='main'.
- **InternVL3-8B**: the pixel-shuffle downsampling stage is the merger; L2 taps its input.
- **M-RoPE fix**: Qwen2.5-VL uses a block M-RoPE layout and Qwen3-VL an interleaved one; under pruning the position cursor must advance by the actual surviving count. A family-scoped fix (bit-degrading to the original at full retention) makes the Qwen2.5-VL compression path well-formed; the block-vs-interleaved contrast is itself a diagnostic of why naive pruning collapses one generation and is tolerated by the other.
- **r2c harness**: HF transformers eager attention, --mrope (family-agnostic correct positions) on all same-scope cells; the Qwen2.5-VL $\times$ pre vllm-mimic positional degeneration (repeated "addCriterion..." outputs $\rightarrow$ pseudo-zero scores) was detected and replaced before any reported number; on Qwen3-VL native vs mimic give 0.575 vs 0.525 (same ordering). The harness and position-layout verification are retained in the audit index.

S9. Per-cell reproducibility index

The filenames below define the anonymous artifact manifest and the audit records required for submission. The archive must be checked against this index before external upload; the tables remain self-contained summaries of the reported cells.

S9.1 InternVL3-8B full-split cells (16/16)

Model OpenGVLab/InternVL3-8B, vLLM 0.19, bf16, eager attention, greedy; $r = \text{TOTAL drop}$ (0.750 $\rightarrow \kappa = 25\%$, 0.875 $\rightarrow \kappa = 12.5\%$); skip = context-overflow cells scored 0. Evidence record: audit index entry for the InternVL3 full-split metrics and per-cell skip accounting.

(Per-cell artifact identifiers are specified by the anonymous manifest. Pre-merger leads on every text-dense benchmark at both budgets — e.g. DocVQA ANLS 0.7284 vs 0.382 at $\kappa = 25\%$; GQA is the near-null object-centric case, 0.5993 vs 0.6031 — replicating the stage law on the pixel-shuffle merger design.)

S9.2 FastV-k3 same-scope cells (8/8)

FastV one-shot attention pruning at LLM layer $K = 3$, iso-budget $r = 0.75$ ($\kappa = 25\%$); HF-transformers 4.57.6 eager harness, official metrics. Official values are rescored outputs in the audit index; n , ptid, skip, engine, and per-cell identifiers are audited there. The harness's raw acc field is not used as the official score.

Model	Benchmark	n	none	pre@25%	post@25%	pre@12.5%	post@12.5%
Qwen3-VL-8B	TextVQA	5000	0.0051	0.0069	0.0059	0.0071	0.0048
Qwen3-VL-8B	DocVQA	5349	— ^a	0.0068	0.0058	0.0065	0.0042
Qwen3-VL-8B	OCR-Bench	1000	0.0135	0.0157	0.0123	0.0151	0.0071
Qwen3-VL-8B	GQA	12578	0.0043	0.0044	0.0045	— ^c	— ^c
Qwen2.5-VL-7B	TextVQA	5000	0.0049	0.0065	0.0070	0.0069	0.0066
Qwen2.5-VL-7B	DocVQA	5349	0.0030	0.0066	0.0068	0.0068	0.0059
Qwen2.5-VL-7B	OCR-Bench	1000	0.0122	0.0158	0.0122	0.0149	0.0079
Qwen2.5-VL-7B	GQA	12578	0.0044	0.0044	0.0044	— ^c	— ^c

Model	Bench	full Δ_{25}	$n = 200$ Δ_{25}	sign
Qwen3-VL	TextVQA	+38.4 pp	+38.3 pp	✓ pre
Qwen3-VL	DocVQA	+24.3 pp	+26.5 pp	✓ pre
Qwen2.5-VL	TextVQA	+26.1 pp	+32.0 pp	✓ pre
Qwen2.5-VL	DocVQA	+11.0 pp	+18.8 pp	✓ pre

Selector	Benchmark	pre	post	Δ
L2	TextVQA	0.598–0.605	0.215–0.222	$\approx +38$ pp
GC-attn	TextVQA	0.553	0.200	+35.3 pp
GC-attn	TextVQA/DocVQA (2n)	—	—	+35.5 / +24.4 pp

Benchmark (metric)	none	pre-final	post	Δ (paired mean $\pm$ SE)	verdict
TextVQA (VQA-acc)	0.858	0.515	0.215	+30.0 pp (+0.300 $\pm$ 0.039)	stage effect
DocVQA (ANLS)	0.976	0.326	0.200	+12.5 pp (+0.125 $\pm$ 0.032)	stage effect
GQA (exact match)	0.605	0.450	0.450	0.0 pp (+0.000 $\pm$ 0.032)	no stage effect

(File names marked _n500 for the Qwen3-VL DocVQA / OCR-Bench cells ran at the $n = 200$ same-scope budget recorded in the JSON n field; OCR-Bench skips are long-OCR context overruns, scored 0.)

S9.3 RBM same-scope cells (5/5)

RBM (pre-merger L2) at iso-budget $r = 0.75$, HF-transformers eager, --mropenative (family-agnostic correct positions), matched to the FastV-k3 cells above. Official values, ptid, skip, and per-cell identifiers are recorded in the audit index.

S9.4 Cascade gate cells (24/24, Qwen3-VL-8B)

Three arms $\times$ two budgets $\times$ four benchmarks; $n = n_{\text{common}}$ (matched by question id); official = acc_official from the paired rescore; ptid = mean post-merger token count (identical across the three arms in every cell). Evidence record: audit index entry for the cascade paired summary and per-arm cells. (McNemar counts for these cells: S4b.)

(Per-cell artifact identifiers are specified by the anonymous manifest. OCR-Bench $n_{\text{common}} = 181$ after 19 long-OCR skips; all other cells $n = 200$. All 24 expected cells are present.)

S10. Qualitative same-image comparison — measured pre- vs post-merger L2 selection

For four representative inputs we render, on the *same* processor-resized image, unit grid, 25% budget and final visual-token count, the original image with its full merge-unit grid, the RBM (pre-merger L2 top- k) kept mask, the post-merger L2 top- k kept mask (same L2 selector, same budget), and a four-state difference map (kept-by-both / RBM-only / post-only / dropped-by-both, encoded by colour *and* hatch/outline so the states remain distinguishable in greyscale). Each panel reports the measured top- k Jaccard index $J = |\text{pre} \cap \text{post}| / |\text{pre} \cup \text{post}|$ and the full-ranking Spearman ρ between the pre- and post-merger L2 orderings (computed from the stored NPZ rank arrays; values cross-checked against the validation report, $|\Delta| < 10^{-5}$). All scores are the target model’s own merger-input / merger-output feature L2 norms (Qwen3-VL-8B-Instruct, enforce_eager, max_pixels = 1.5M); no CLIP, SigLIP, attention map, or fabricated saliency is used. The FastV layer-2 panel is intentionally absent (no user-supplied question exists for these images; a query-conditioned mask cannot be measured without one) and is recorded as fastv_status: skipped_missing_question.

The consistently low Jaccard (0.25–0.35) and moderate Spearman (0.30–0.48) across the four shown inputs are the image-level signature of the body paper’s mechanism result: the learned 2×2 merger *reshuffles* the unit ranking, so selecting before versus after it on the same image with the same selector and budget keeps largely

different units. The four-state maps show that the disagreement is spatially structured (contiguous RBM-only and post-only regions), not uniform noise, and that the “kept-by-both” intersection is a sparse minority of the budget.

Packaging note: this source package contains the rendered summaries above, not the run-level anonymous artifact indexed in S9. That archive must be restored, sanitized, checksummed, and verified before submission. This supervisor-review draft identifies Zhengxing Yan but includes no affiliation or repository identifier; a future anonymous submission requires a separate anonymization pass.

Body location	Cells	Artifact record	Audit summary
Main full-split Qwen blocks	26/26	audit: main-qwen	full-split metrics and skip accounting
Main full-split InternVL3 block	16/16	audit: main-internvl3	full-split metrics and skip accounting
Regime-map FastV-k3	8	audit: regime-fastv-k3	same-scope rescored cells
Regime-map RBM same-scope	5	audit: regime-rbm	same-scope rescored cells
Cascade gate	24/24	audit: cascade	paired gate statistics
Efficiency study	7 configs	audit: efficiency	throughput and latency records
§ 5 kept-set Jaccard / swap	8	audit: mechanism-swap	kept-set and ordering audit
§ 5.1–5.2 M1/M2	$n = 64 \times 3$	audit: mechanism-m1-m2	rank and edge-energy measurements
§ 5.4 retention curve	$n = 200 \times 4$	audit: retention-curve	depth ablation records
§ 6 (a)–(f) prespecified extensions	—	audit: negative-results	gates and extension records
Learned-scorer probe	$n = 200 \times 2$	audit: learned-scorer	model, held-out splits, predictions, and official rescoring
§ 8 600k reversal	3	audit: low-cap-reversal	configuration-boundary records
Engine fairness	$r=0$ 8/8; none 16/16	audit: engine-parity	zero-drop and none-cell checks
PyramidDrop iso-25% col-lapse	$n = 500$	audit: pyramiddrop-iso25	iso-budget baseline record


cell	mode	r	bench	n	official	metric	ptid	skip	JSON
none_docvqa	none	0.000	docvqa	5349	0.9221	ANLS	3230.8	0	internvl3_none_docvqa_r0.000_full.json
none_gqa	none	0.000	gqa	12578	0.6293	acc-official	2557.2	0	internvl3_none_gqa_r0.000_full.json
none_ocrbench	none	0.000	ocrbench	1000	0.852	Final/1000=852	2040.1	0	internvl3_none_ocrbench_r0.000_full.json
none_textvqa	none	0.000	textvqa	5000	0.8338	VQA-acc	2562.9	0	internvl3_none_textvqa_r0.000_full.json
pre_docvqa_r0.750	pre	0.750	docvqa	5349	0.7284	ANLS	859.5	0	internvl3_pre_docvqa_r0.750_full.json
pre_docvqa_r0.875	pre	0.875	docvqa	5349	0.5054	ANLS	464.3	0	internvl3_pre_docvqa_r0.875_full.json
pre_gqa_r0.750	pre	0.750	gqa	12578	0.5993	acc-official	690.1	0	internvl3_pre_gqa_r0.750_full.json
pre_ocrbench_r0.750	pre	0.750	ocrbench	1000	0.753	Final/1000=753	555.4	0	internvl3_pre_ocrbench_r0.750_full.json
pre_textvqa_r0.750	pre	0.750	textvqa	5000	0.789	VQA-acc	690.5	0	internvl3_pre_textvqa_r0.750_full.json
pre_textvqa_r0.875	pre	0.875	textvqa	5000	0.723	VQA-acc	378.4	0	internvl3_pre_textvqa_r0.875_full.json
post_docvqa_r0.750	post	0.750	docvqa	5349	0.382	ANLS	859.5	0	internvl3_post_docvqa_r0.750_full.json
post_docvqa_r0.875	post	0.875	docvqa	5349	0.2451	ANLS	464.3	0	internvl3_post_docvqa_r0.875_full.json
post_gqa_r0.750	post	0.750	gqa	12578	0.6031	acc-official	690.1	0	internvl3_post_gqa_r0.750_full.json
post_ocrbench_r0.750	post	0.750	ocrbench	1000	0.321	Final/1000=321	555.4	0	internvl3_post_ocrbench_r0.750_full.json
post_textvqa_r0.750	post	0.750	textvqa	5000	0.4148	VQA-acc	690.5	0	internvl3_post_textvqa_r0.750_full.json
post_textvqa_r0.875	post	0.875	textvqa	5000	0.3064	VQA-acc	378.4	0	internvl3_post_textvqa_r0.875_full.json

cell	model	bench	n	official	ptid	skip	engine	JSON
r2b_qwen3vl_fastv_k3_textvqa_r0.75_full5000	Qwen3-VL	textvqa	5000	0.7771	213.7	5	hf eager	r2b_qwen3vl_fastv_k3_textvqa_r0.75_full5000.json
r2b_qwen3vl_fastv_k3_gqa_r0.75_full12578	Qwen3-VL	gqa	12578	0.5376	96.8	0	hf eager	r2b_qwen3vl_fastv_k3_gqa_r0.75_full12578.json
r2b_qwen3vl_fastv_k3_docvqa_r0.75_n500	Qwen3-VL	docvqa	200	0.5863	176.0	0	hf eager	r2b_qwen3vl_fastv_k3_docvqa_r0.75_n500.json
r2b_qwen3vl_fastv_k3_ocrbench_r0.75_n500	Qwen3-VL	ocrbench	200	0.415	114.8	19	hf eager	r2b_qwen3vl_fastv_k3_ocrbench_r0.75_n500.json
r2b_qwen2vl_fastv_k3_textvqa_r0.75_n200	Qwen2.5-VL	textvqa	200	0.7467	282.1	0	hf eager	r2b_qwen2vl_fastv_k3_textvqa_r0.75_n200.json
r2b_qwen2vl_fastv_k3_gqa_r0.75_n200	Qwen2.5-VL	gqa	200	0.475	126.8	0	hf eager	r2b_qwen2vl_fastv_k3_gqa_r0.75_n200.json
r2b_qwen2vl_fastv_k3_docvqa_r0.75_n200	Qwen2.5-VL	docvqa	200	0.4852	231.4	0	hf eager	r2b_qwen2vl_fastv_k3_docvqa_r0.75_n200.json
r2b_qwen2vl_fastv_k3_ocrbench_r0.75_n200	Qwen2.5-VL	ocrbench	200	0.285	143.6	20	hf eager	r2b_qwen2vl_fastv_k3_ocrbench_r0.75_n200.json

cell	model	bench	n	official	ptid	skip	engine	JSON
r2c_qwen2vl_pre_r0.75_textvqa_n200	Qwen2.5-VL	textvqa	200	0.6683	282.1	0	hf eager	r2c_qwen2vl_pre_r0.75_textvqa_n200.json
r2c_qwen2vl_pre_r0.75_gqa_n200	Qwen2.5-VL	gqa	200	0.520	126.8	0	hf eager	r2c_qwen2vl_pre_r0.75_gqa_n200.json
r2c_qwen2vl_pre_r0.75_docvqa_n200	Qwen2.5-VL	docvqa	200	0.5062	231.4	0	hf eager	r2c_qwen2vl_pre_r0.75_docvqa_n200.json
r2c_qwen2vl_pre_r0.75_ocrbench_n200	Qwen2.5-VL	ocrbench	200	0.370	143.6	20	hf eager	r2c_qwen2vl_pre_r0.75_ocrbench_n200.json
r2c_qwen3vl_pre_r0.75_ocrbench_n200	Qwen3-VL	ocrbench	200	0.575	114.8	19	hf eager	r2c_qwen3vl_pre_r0.75_ocrbench_n200.json

arm	κ	bench	n	official	ptid	JSON
pre-alone (RBM)	25%	textvqa	200	0.5967	212.8	gate_pre25_textvqa.json
fastv-alone	25%	textvqa	200	0.6800	212.8	gate_fst25_textvqa.json
cascade	25%	textvqa	200	0.5950	212.8	gate_cas25_textvqa.json
pre-alone (RBM)	25%	docvqa	200	0.4239	176.0	gate_pre25_docvqa.json
fastv-alone	25%	docvqa	200	0.5183	176.0	gate_fst25_docvqa.json
cascade	25%	docvqa	200	0.4779	176.0	gate_cas25_docvqa.json
pre-alone (RBM)	25%	ocrbench	181	0.5801	114.8	gate_pre25_ocrbench.json
fastv-alone	25%	ocrbench	181	0.1713	114.8	gate_fst25_ocrbench.json
cascade	25%	ocrbench	181	0.3978	114.8	gate_cas25_ocrbench.json
pre-alone (RBM)	25%	gqa	200	0.4150	95.4	gate_pre25_gqa.json
fastv-alone	25%	gqa	200	0.4900	95.4	gate_fst25_gqa.json
cascade	25%	gqa	200	0.4050	95.4	gate_cas25_gqa.json
pre-alone (RBM)	12.5%	textvqa	200	0.5000	120.7	gate_pre12_textvqa.json
fastv-alone	12.5%	textvqa	200	0.4450	120.7	gate_fst12_textvqa.json
cascade	12.5%	textvqa	200	0.4050	120.7	gate_cas12_textvqa.json
pre-alone (RBM)	12.5%	docvqa	200	0.2980	103.7	gate_pre12_docvqa.json
fastv-alone	12.5%	docvqa	200	0.2779	103.7	gate_fst12_docvqa.json
cascade	12.5%	docvqa	200	0.2070	103.7	gate_cas12_docvqa.json
pre-alone (RBM)	12.5%	ocrbench	181	0.3591	68.7	gate_pre12_ocrbench.json
fastv-alone	12.5%	ocrbench	181	0.0884	68.7	gate_fst12_ocrbench.json
cascade	12.5%	ocrbench	181	0.1713	68.7	gate_cas12_ocrbench.json
pre-alone (RBM)	12.5%	gqa	200	0.3950	62.3	gate_pre12_gqa.json
fastv-alone	12.5%	gqa	200	0.4550	62.3	gate_fst12_gqa.json
cascade	12.5%	gqa	200	0.3500	62.3	gate_cas12_gqa.json

Same-image iso-budget comparison - df7282e1 (Qwen3-VL-8B measured merger L2)



$N=1144$ $k=286$ (25% kept) top-k Jaccard=0.2599 Spearman $\rho(\text{pre}, \text{post ranks})=0.3141$ unit grid $26 \times 44 = \text{grid_thw}/2$ Qwen3-VL-8B measured merger L2


FastV: SKIPPED (no user question; inputs/manifest.json absent)

(a) full unit grid outline (light). (b,c) SAME amber L2 selector & 25% budget; (c) hatched = post (distinguishable without color).
(d) four states kept-by-both / RBM-only / post-only / dropped-by-both, encoded by color + hatch/outline (legend).

■ kept by both (pre & post) ■ RBM-only (pre & ~post) ■ post-only (~pre & post) ■ dropped by both

Figure 6: Sample df7282e1 ($N = 1144$, $k = 286$, $J = 0.260$, $\rho = 0.314$). Sparse scene; the four-state map shows RBM-only (amber) concentrated on object boundaries while post-only (hatched) scatters over flat regions the merger smoothed.

Same-image iso-budget comparison - 69 (Qwen3-VL-8B measured merger L2)



$N=6820$ $k=1705$ (25% kept) top-k Jaccard=0.3478 Spearman $\rho(\text{pre}, \text{post ranks})=0.4774$ unit grid $110 \times 62 = \text{grid_thw}/2$ Qwen3-VL-8B measured merger L2


FastV: SKIPPED (no user question; inputs/manifest.json absent)

(a) full unit grid outline (light). (b,c) SAME amber L2 selector & 25% budget; (c) hatched = post (distinguishable without color).
 (d) four states kept-by-both / RBM-only / post-only / dropped-by-both, encoded by color + hatch/outline (legend).

■ kept by both (pre & post) ■ RBM-only (pre & ~post) ■ post-only (~pre & post) ■ dropped by both

Figure 7: Sample *img69* ($N=6820$, $k=1705$, $J=0.348$, $\rho=0.477$). Highest overlap of the four; the merger preserves more of the pre-ranking here, yet two-thirds of the kept set still differs.

Same-image iso-budget comparison - 72 (Qwen3-VL-8B measured merger L2)



$N=16224$ $k=4056$ (25% kept) $\text{top-k Jaccard}=0.2482$ $\text{Spearman } \rho(\text{pre, post ranks})=0.2999$ $\text{unit grid } 104 \times 156 = \text{grid_thw}/2$ Qwen3-VL-8B measured merger L2

FastV: SKIPPED (no user question; inputs/manifest.json absent)


(a) full unit grid outline (light). (b,c) SAME amber L2 selector & 25% budget; (c) hatched = post (distinguishable without color).
(d) four states kept-by-both / RBM-only / post-only / dropped-by-both, encoded by color + hatch/outline (legend).

■ kept by both (pre & post) ■ RBM-only (pre & ~post) ■ post-only (~pre & post) ■ dropped by both


Figure 8: Sample img72 ($N=16224$, $k=4056$, $J=0.248$, $\rho=0.300$). Textured outdoor scene with a wall edge; the lowest overlap – the merger’s lossy aggregation most strongly reorders the high-frequency edge vs. texture units.

Same-image iso-budget comparison - 74 (Qwen3-VL-8B measured merger L2)


a original




b RBM (pre L2)



c post L2



d pre vs post



$N=8160$ $k=2040$ (25% kept) top-k Jaccard=0.3377 Spearman $\rho(\text{pre}, \text{post ranks})=0.3963$ unit grid $68 \times 120 = \text{grid_thw}/2$ Qwen3-VL-8B measured merger L2

FastV: SKIPPED (no user question; inputs/manifest.json absent)

(a) full unit grid outline (light). (b,c) SAME amber L2 selector & 25% budget; (c) hatched = post (distinguishable without color).

(d) four states kept-by-both / RBM-only / post-only / dropped-by-both, encoded by color + hatch/outline (legend).

■ kept by both (pre & post) ■ RBM-only (pre & ~post) ■ post-only (~pre & post) ■ dropped by both

Figure 9: Sample img74 ($N=8160$, $k=2040$, $J=0.338$, $\rho=0.396$). Waterscape with horizon; RBM-only units track the water surface detail that the post-merger L2 demotes in favour of merged sky blocks.